

ActiveClean: Progressive Data Cleaning For Advanced Analytics

ABSTRACT

There have been numerous proposals to extend data management systems and frameworks with support for machine learning. A perennial challenge in these systems is the presence of dirty data – in the form of missing, duplicate, incorrect or inconsistent values – which can bias the learning algorithms, leading to incorrect or error-prone predictions. Although this can be ameliorated by cleaning the input dataset using existing data cleaning techniques, the cost of doing so (e.g., manually finding and cleaning all dirtiness) is practically and economically infeasible. In addition, existing learning algorithms are not suited for training on partially cleaned datasets, which can lead to dramatically incorrect models. In response, we have designed and implemented the ActiveClean system that integrates incremental data cleaning with machine learning algorithms. ActiveClean iteratively updates a machine learning model as data records are individually (or in batches) updated. We additionally study numerous database and ML-oriented optimizations to speed the end-to-end process. Our evaluation on four real-world datasets suggest that our methodology can learn more accurate models by cleaning significantly less data records than existing alternatives such as uniform sampling and active learning, thus bridging the gap between the reality of dirty data sources and final data products such as predictive models.

1. INTRODUCTION

There have been several recent proposals to extend existing data management systems with advanced analytics that go beyond SQL aggregates (e.g., MLBase [25], MADLIB [20], TupleWare [10]). While these extensions can dramatically reduce the required effort in developing Machine Learning and other predictive applications, pre-processing or “cleaning” data is still a serious bottleneck [22]. Data are susceptible to various forms of corruption such as missing, incorrect, or inconsistent values. Industrial surveys have established that dirty data are prevalent [38], and there is a growing industry around cleaning data [1].

Machine Learning is known to be sensitive to data quality [44], and improper handling of dirty data can lead to unreliable and error-prone predictions. Supervised Machine Learning relies on correlating features with labels, and systematically corrupted data can lead to confounding correlations. Consider a music recommender system in which due to a software bug, all users from Europe have an incorrect age attribute defaulted to “18-24”. A recommendation model trained on this data may spuriously “learn” a correlation relationship between the “18-24” age group and music liked by European users. While there is an extensive literature on robust Machine Learning, this work largely focuses on resilience to atypical outliers (i.e., age “150”) and not spurious correlations.

Such problems can be ameliorated by a variety of data cleaning techniques (see Rahm and Do [34] for a survey). Recently, several scalable data cleaning frameworks have been proposed including




Figure 1: ActiveClean presents an architecture where data cleaning is integrated with model training. The user specifies a data cleaning procedure, and we provide a framework for sampling, model update, and feedback through estimation.

BigDancing [23], SampleClean [4], and Katara [9]. Unfortunately, even with recent research results, data cleaning can be time consuming and economically infeasible [43]. Programming data transformations to fix all problems manifest in the data requires substantial developer effort [22]. When scripted transformations do not suffice, crowdsourcing is a popular alternative with success in missing value filling and entity resolution ([4,9,18,32]). However, crowdsourcing comes at the cost of additional latency and significant monetary expense.

Thus, for many corrupted datasets, the analyst can only practically clean $k \ll N$ records. Strict budgets pose a significant challenge to Machine Learning. Machine Learning applications can require a large amount of training data to learn correlations between possibly hundreds of features and labels. A small sample of clean data may not be able to support a viable Machine Learning model. Likewise, training on partially cleaned data, where cleaned data is mixed with dirty data, can lead to dramatically incorrect models (Figure 2). This leads to the central question that we explore in this work, namely, when and how to apply data cleaning in a Machine Learning model training pipeline to avoid the severe dependence on sample size.

In this paper, we propose ActiveClean, a framework that integrates existing data cleaning solutions with Machine Learning model training in a way that ensures reliable intermediate results. Suppose an analyst has a large dirty dataset, a data cleaning technique, a budget of cleaning k records, and a model she wishes to train. ActiveClean uses the model to select the most valuable samples of data to clean, and incrementally updates the model after cleaning. The cleaning is iterative, applying the data cleaning procedure to samples and feeding the results back to adaptively select the next sample. ActiveClean can integrate user specified dirty data detection rules (such as constraints) into the framework to further guide our sampling. ActiveClean supports a broad class of models, called convex-loss models, and a variety of data cleaning techniques (e.g., constraint-based, value filling, and entity resolution). The key insight is that the iterative data cleaning and update process can be analyzed as a form of Stochastic Optimization [7] allowing us to leverage convex-analysis theory to develop optimized

sampling and estimation methods.

As analytics frameworks increasingly support Machine Learning, tools that encourage appropriate methodologies are required. The incremental updates from ActiveClean give provable bounds on intermediate results and these results are also viable models that can be used for prediction. Specifically, our contributions are:

- We propose a technique that updates a dirty model from clean data based on gradient descent.
- We derive an optimal sampling distribution for estimating the gradient descent step, which while practically incomputable, can be approximated by maintaining an estimate of the impact of data cleaning on a record.
- The estimation technique applies linearization to decouple changes in different features allowing us to integrate knowledge about what is wrong with the data (i.e. analyst-specified detection rules) to better estimate error impact.
- We show that batching together data that we expect to be clean can lead to significant improvements especially when the sample size is small.
- We evaluate ActiveClean on real and synthetic datasets to show that our approach gives a sharper cost-accuracy trade-off than random sampling and Active Learning.

2. PROBLEM SETUP

2.1 Examples of Systematic Corruption

We first introduce a few examples of systematic corruption in publically available data.

Dollars for Docs: ProPublica has collected a dataset of corporate contributions to doctors for analysis of whether these contributions affect research. On their website [2], they list numerous types of corruption that had to be cleaned before publishing the data. The most significant contributions were made by large companies whose names were more often inconsistently represented in the data e.g. “Pfizer Inc.”, “Pfizer Incorporated”, “Pfizer”. If this corruption is not cleaned, contributions from larger companies might be underestimated.

Country Imports: The monetary value of a country’s imports are an important metric for understanding purchasing power, balance of payments, and economic strength. Import statistics available on publically curated datasets such as Wikipedia [3] are up-to-date for large countries but are 8-10 years behind for smaller countries. Smaller countries often experience more rapid growth (or decline) than larger countries making this systematic corruption problematic.

2.2 Motivating Analysis Scenario

Consider the following motivating scenario. An analyst is interested in understanding the correlation between neurological data (EEG data), patient predispositions, and the onset of seizures.

EXAMPLE 1. *An analyst is given a 15-dimensional feature vector $x \in \mathbb{R}^{15}$ of sampled from a few seconds of electrical signals and three binary features indicating risk factors including family history, age, and gender. The classification is a label $y \in \{0, 1\}$ of whether the patient is having a seizure. The analyst wants to train a linear regression model to predict seizures.*

Medical records are often dirty due to digitization and data integration [1]. After model training, the analyst is informed of corruption in the patient data. With some preliminary data exploration, she discovers that older patients whose records were manually digitized are more likely to have missing values. In addition, occurrences of certain types of seizures are mislabeled based on a changing medical definition. The procedure to determine which patients’

data are corrupted requires querying source medical records which are in a semi-structured format.

Cleaning With a Budget: Let us suppose our analyst wants evaluate whether the problems affect her model, and she wants to do this without having to manually validate every record. The first solution to this problem is to clean a sample of data, write this sample back, and then retrain the model on the partially cleaned data. However, this solution can give unreliable results since we are mixing dirty and clean data. In Figure 2, we illustrate the dangers of this approach on a very simple dirty dataset and model. Suppose, we are trying to find a best fit line for two variables. One of the variables is systematically corrupted with a translation in the x-axis (Figure 2a). We mark the dirty data in brown and the clean data in orange, and show their respective best fit lines in blue. If we clean two of the data points (Figure 2b), and then find the best fit line, the partially cleaned model is dramatically incorrect. This is a well-known phenomenon called Simpson’s paradox, where mixtures of different populations of data can result in spurious relationships [37]. The result is surprising, namely, Machine Learning on partially cleaned data can lead to misleading results.

The next solution is avoiding the dirty data altogether. She can take a sample of data, apply data cleaning, and train a model only on that sample of data. This is similar to SampleClean [43]; proposed to approximate the results to aggregate queries by applying them to a clean sample of data. Machine Learning models are highly sensitive to sample size when high dimensional. The result is that the models will not be useful (i.e., cannot accurately predict test data) when the sample size is too small (Figure 2c).




Figure 2: (a) We present a dataset with a systematic corruption (translation) in one variable where the dirty data is in brown and the clean data is in yellow. If we draw the best fit line (i.e., linear regression), we see that the systematic corruption affects the line. (b) However, after cleaning two data points, applying the same model to partially clean data results in a dramatically incorrect answer. (c) Likewise, small sample sizes can result in similarly incorrect models.

ActiveClean avoids both pitfalls, Simpson’s paradox and sample size dependence. In Section 5, we show how we do this with iterative gradient steps (i.e., incrementally moving the line based on the clean data). This takes advantage of the dirty data as well as the clean data, but still have provable properties about the intermediate results. In Figure 3, we illustrate our ideal tradeoff space of sampling and data cleaning. At two extremes we have no cleaning (just using the dirty data) and full cleaning. ActiveClean is optimized for convergence for smaller sample sizes than a uniform sampling approach.

2.3 Preliminaries

We focus on a class of Supervised Learning problems called Empirical Risk Minimization (ERM) (see Friedman, Hastie, and Tibshirani [17] for an introduction). In ERM, the goal is to learn a set of model parameters θ from training examples. We start with a set of training examples $\{(x_i, y_i)\}_{i=1}^N$ on which we minimize a loss function ϕ (a penalty for getting the prediction wrong) at each point




Figure 3: Our goal with ActiveClean is designed to converge to an accurate model with fewer cleaned records than a uniform sampling approach (SampleClean).

parametrized by θ . A very important class of problems is when the loss function is convex in θ . Convex problems are amenable to a variety of different optimization methodologies and have strong theoretical guarantees on convergence rates. This class of problems includes all generalized linear models (including linear and logistic regression), and all variants of support vector machines. Formally,

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^N \phi(x_i, y_i, \theta)$$

For example, in a linear regression ϕ is:

$$\phi(x_i, y_i, \theta) = \|\theta^T x_i - y_i\|_2^2$$

Typically, a *regularization* term $r(\theta)$ is added to this problem. $r(\theta)$ penalizes high or low values of feature weights in θ to avoid over-fitting to noise in the training examples.

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^N \phi(x_i, y_i, \theta) + r(\theta)$$

For example, a popular variant of linear regression is called LASSO which is:

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^N \|\theta^T x_i - y_i\|_2^2 + \lambda \cdot \|\theta\|_1$$

By applying the L1 regularization term, if one of the features is particularly noisy, and does not add predictive value, it will get excluded. The loss function specifies a problem independent of the optimization used to calculate the optimal parameter θ^* .

2.3.1 Types of Errors

When discussing “errors” in Machine Learning models, we refer to the following metrics:

Model Error. Let θ be the model trained on the dirty data, and let θ^* be the model trained on the same data if it were cleaned. Then the model error is defined as $\|\theta - \theta^*\|$.

Testing Error. Let θ be the model trained on the dirty data, and let θ^* be the model trained on the same data if it were cleaned. Let $T(\theta)$ be the out-of-sample testing error when the dirty model is applied to the clean data, and $T(\theta^*)$ be the test error when the clean model is applied to the clean data. The testing error is defined as $T(\theta^*) - T(\theta)$.

The *error* in the model and in testing is caused by two underlying problems:

Data Error. Error caused by dirtiness.

Sampling Error. Error introduced by sampling i.e., the difference between a model trained on $p\%$ of the data and 100% of the data.

In this paper, when we use the term *error*, we are referring to **model error**. We will be explicit when we are using the word error to describe problems regarding dirtiness in the underlying data (e.g. “data corruption”).

2.4 ActiveClean Problem

The basic question that we explore in this work is how to update a model trained on dirty data with data cleaning. Formally,

PROBLEM 1 (ACTIVECLEAN PROBLEM). Let R be a dirty relation and each row $r \in R$ is turned into a feature vector and label tuple $F(r) \mapsto (x, y)$. We are given a convex regularized loss model parametrized by θ trained on the set of features and labels $\{(x, y)\}$. The user specifies a data repair step which cleans a record $C(r) \mapsto r_{clean}$. Optionally, the user can provide a set of dirty data detection conditions that select a set of candidate dirty data R_{dirty} . With a budget of applying cleaning i.e., $C(\cdot)$, only k times, return an estimate $\hat{\theta}$ of the clean model.

Our goal is to use information from the model to inform data cleaning on samples, and use the information from clean samples to update the model. The tight feedback loop between model training and data cleaning poses several new algorithmic challenges. If we are not careful about how we update the model, we can end up with the Simpsons paradox problem described before. Data cleaning and model training happen at very different time scales, we have to carefully budget our effort to ensure that optimizations address rate-determining steps in the workflow.

This framework is optimized around problems with expensive data cleaning. As the complexity of dirtiness increases, so does the amount of computation and human-effort needed to resolve them. The consequence is that automated repair is either not scalable or relies on greedy repairs that can be unreliable. Increasingly, human effort is seen as valuable in data cleaning [18,32,42,43]. When humans are involved, per record latencies for data repair are orders of magnitude larger than the CPU time needed for model training. We can compare recent results in data cleaning to a model training framework like CoCoA implemented on Spark [21]. Per record, BigDancing, a highly optimized automated Spark-based data cleaning system is 15.5x slower than CoCoA¹. Crowd based techniques like CrowdFill [32] and CrowDER [42] are over 100,000x slower per record.

Here is an example application of ActiveClean with our running example:

EXAMPLE 2. The analyst first trains her linear model on the dirty data ignoring the effects of the errors returning a model $\theta^{(d)}$. She decides that she has a budget of cleaning 40 records, and decides to clean the 40 records in batches of 10 (set based on how fast she can clean the data, and how often she wants to see an updated result). She initializes ActiveClean with $\theta^{(d)}$. ActiveClean samples an initial batch of 10 records. She manually cleans those records by using the medical record source data. After each batch, the model is updated with the most recent cleaning results $\theta^{(t)}$. The model improves after each iteration. After $t = 4$ of cleaning, the analyst has an accurate model trained with 40 cleaned records but still utilizes the entire dirty data.

3. SYSTEM ARCHITECTURE

We describe the ActiveClean architecture and the basic algorithmic framework. We will address the individual components in the subsequent sections.

3.1 Overview

We revisit the architecture diagram (Figure 1) in more detail. The first step of the system is *initialization*. In this step, the user provides a relation R , a data cleaning technique $C(\cdot)$, and a dirty model $\theta^{(d)}$ trained on the dirty dataset. Optionally, the user can provide a dirty data detector $D(\cdot)$ which selects the set of likely

¹For CoCoA to reach a precision of 1e-3

corrupted records from R . This is possible for some types of data errors such as constraint violations and missing values. If one is not provided, we start by treating all of the data is dirty and try to learn a detector as we clean more data. We are given a *featurization* function $F(\cdot)$ that maps each record in R to a feature vector and a label. At initialization, there are two hyperparameters to set, the cleaning budget k and the batch size b (the number of iterations is $t = \frac{k}{b}$). We discuss how to set b and the tradeoffs in setting a larger or smaller b in Section .

After initialization, ActiveClean begins the cleaning and model update iterations. The *sampler* selects a sample of dirty data based on the batch size. At this step, ActiveClean can use the detector D to narrow the sample to select only dirty data. Once a sample is selected, we apply the user-specified data cleaning operation to clean the sample (the *cleaner*). Then, after the sample is cleaned, the dirty model is updated by the *updater*.

The next two steps in the architecture are feedback steps where we update the sampling for the next iteration. The *estimator* estimates the change on the model if a dirty record is cleaned based on prior cleaned data. This information can be used to guide sampling towards more valuable records. After estimation, we also update our detector $D(\cdot)$ based on the data that we have cleaned. After all of the iterations are complete, the system returns the updated model.

To summarize the architecture in pseudocode:

```

1. Init(dirty_data, dirty_model, batch, iter)
2. For each i in {1, ..., T}
    (a) dirty_sample = Sampler(dirty_data, sample_prob,
        detector, batch)
    (b) clean_sample = Cleaner(dirty_sample)
    (c) current_model = Updater(current_model,
        sample_prob, clean_sample)
    (d) sample_prob = Estimator(dirty_data,
        cleaned_data, detector)
    (e) detector = DetectorUpdater(detector)
3. Output: current_model

```

3.2 Challenges and Formalization

We highlight the important components and formalize the research questions that we explore in this paper.

Featurization. Given a relation R with a set of attributes A . There is a featurization function F which maps every row in \mathcal{R} to a d dimensional feature vector and a l dimensional label tuple:

$$F(r \in R) \rightarrow (\mathbb{R}^l, \mathbb{R}^d)$$

The result of the featurization are the data matrices X and Y .

$$F(R) \rightarrow (X, Y)$$

Detector (Section 4). The first challenge in ActiveClean is dirty data detection. This is key part of our sampling framework. In this step, we select a candidate set of dirty records $R_{dirty} \subseteq R$. We will discuss two techniques to do this: (1) an a priori case, and (2) an adaptive case. In the a priori case, we know which data is dirty in advance. In the adaptive case, we train a classifier based on data that we have already cleaned to select the dirty data.

Sampler (Section 6). We take a sample of the records $S_{dirty} \subseteq R_{dirty}$. This is a non-uniform sample where each record r has a sampling probability p_r . We will show that we can derive the optimal sampling distribution theoretically, however, it will depend on knowing the cleaned data before we actually clean. This is why we introduce an estimation step to approximate this cleaned value.

Updater (Section 5). Then, we update the model $\theta^{(i)}$ based on the newly cleaned data $F(S_{clean})$. The result is the updated model $\theta^{(i+1)}$. We will present the model update before we derive the

optimal sampling distribution. This is because that distribution will naturally fall out of our update procedure.

Estimator (Section 7): The estimator approximates the optimal distribution derived in the Sample step. Based on the change between $F(S_{clean})$ and $F(S_{dirty})$, we direct the next iteration of sampling to select points that will have changes most valuable to the next model update.

4. DETECTION

An important part of ActiveClean is dirty data detection. When we sample data to clean, we want to ensure that the data that we are sampling is actually dirty. Therefore, before we can sample, we have to detect whether the record is likely to be dirty.

4.1 Goals

The detection component needs to give us two important pieces of information about a record: (1) whether the record is dirty, and (2) if it is dirty what is wrong with the record. From (1), we can select a subset of dirty records to sample at each batch. (2) is important so we can estimate how valuable cleaning will be for this record. There are two cases that we explore in ActiveClean: *a priori* and adaptive. In the a priori case, we recognize that for many data cleaning methodologies, we can efficiently select the set of dirty records without repair. In the adaptive case, we relax this assumption, and explore how we can learn which records are dirty and clean with a classifier.

4.2 A Priori Case

For many types of dirtiness such as missing attribute values and constraint violations, it is possible to efficiently enumerate a set of corrupted records and enumerate what is wrong with these records.

DEFINITION 1 (A PRIORI DETECTION). Let r be a record in R . An a priori detector is a detector that returns a Boolean of whether the record is dirty and a set of columns e_r that are dirty.

$$D(r) = (\{0, 1\}, e_r)$$

From the set of columns that are dirty, we can find the corresponding features that are dirty f_r .

We highlight example use cases of this definition using data cleaning methodologies proposed in the literature.

Constraint-based Repair: One model for handling errors in database declaring a set of constraints on a database. Data are cleaned iteratively until the constraints are satisfied [15,23,45].

Detection. Let Σ be a set of constraints on the relation \mathcal{R} . In the detection step, we select a subset of records $\mathcal{R}_{dirty} \subseteq \mathcal{R}$ that violate at least one constraint. The set e_r is the set of columns for each record which have a constraint violation.

EXAMPLE 3. Consider our EEG data running example, for an example of constraint-based cleaning. We can add a constraint that one of the electrical signals cannot be greater than 4V. For all records whose value is above 4V, we would select them in the error detection step.

Entity Resolution: Another common data cleaning task is Entity Resolution [18,24,42]. Entity Resolution is the problem of standardizing attributes that represent the same real world entity. A common pattern in Entity Resolution is to split up the operation into two steps: blocking and matching. In blocking, attributes that should be the same are coarsely grouped together. In matching, those coarse groups are resolved to a set of distinct entities.

Detection. This is the matching step. Let S be a similarity function that takes two records and returns a value in $[0, 1]$ (1 most

similar and 0 least similar). For some threshold t , S defines a similarity relationship between two attributes $r(a)$ and $r'(a)$:

$$r(a) \approx r'(a) : S(r(a), r'(a)) \geq t$$

In the detection step, R_{dirty} is the set of records that have at least one other record in the relation that satisfy $r(a) \approx r'(a)$. The set e_r is the attributes of r that have entity resolution problems.

EXAMPLE 4. An example of an Entity Resolution problem is where the patient’s gender is inconsistently represented (e.g., “Male”, “M”, “Man”). We can define a similarity relationship (first char = ‘M’) and select all records that satisfy this condition.

4.3 Adaptive Detection

Next, we explore how we can handle the case where the detection is not known in advance. Based on what we have already cleaned, we can learn a detection classifier (note that this “learning” is distinct from the “learning” at the end of the pipeline). The challenge in formulating this problem is that we not only need to know whether or not a record is dirty, but also how it is dirty (e.g. e_r in the a priori case). We try to formulate a technique that generalizes what we did before. Instead of assuming that we know which features are corrupted, let us say that we know that there are u classes of data corruption. These classes are corruption categories that do not necessarily align with features, but every records is classified with at most one category. For example, suppose we have outliers and missing values, there are three classes of corruption: outliers, missing values, and both. As the analyst cleans data, she tags dirty data with one of the u classes. Then, the detection problem reduces to a multiclass classification problem. To address this problem, we can use any multiclass classifier, and we use an all-versus-one SVM in our experiments. Since this classifier is internal to our system, it does not have to be a convex model (i.e., it can be a Decision Tree or Random Forest).

When an example (x, y) is cleaned, the repair step also has to provide a label to which of the clean, 1, 2, ..., u classes it belongs. It is possible that u increases each iteration as more types of dirtiness are discovered. Thus, after cleaning k records, we have a dataset of k records labeled with $u + 1$ classes (including one for “not dirty”).

DEFINITION 2 (ADAPTIVE CASE). To select R_{dirty} , we select the set of records for which κ gives a positive error classification (i.e., one of the u error classes). After each sample of data is cleaned, the classifier κ is retrained. So the result is:

$$D(r) = (\{1, 0\}, \{1, \dots, u + 1\})$$

We highlight an example of adaptive detection using an interactive data cleaning tool such as OpenRefine [39].

Interactive Data Cleaning.

EXAMPLE 5. OpenRefine is a spreadsheet-based tool that allows users to explore and transform data. However, it is limited to clean data that can fit in memory on a single personal computer. Thus, it is beneficial to sample data. Since the cleaning operations are coupled with data exploration, we do not know what is dirty in advance (the analyst may discover new errors as she cleans).

Suppose our analyst wants to use OpenRefine to clean a very large dataset with ActiveClean. She takes a sample of data from the entire dataset. Every time she makes a batch data transformation, we find the set of records that have changed. Using the set of distinct transformations, we can infer the corruption classes. These records become positive training examples for a classifier to guide future samples.

5. UPDATE

Before we discuss sampling, we will first discuss how to update a model. We will assume that our detection step in the previous section has selected a set of candidate records R_{dirty} and our sampling

step has sampled from this set of candidate records. We will show that this model update procedure can be interpreted as a Stochastic Gradient Descent (SGD) algorithm, which gives us a theoretical framework to analyze convergence and bound the error at each step.

5.1 Update Problem

The challenge is that we want a model update technique that takes advantage of the entire dirty data and a small sample of cleaned data. However, we want to avoid the Simpson’s paradox problem discussed previously. Let us describe the abstract problem. Suppose we have a random sample of data $S_{dirty} \subseteq R_{dirty}$ from the results of the detection in the previous section. We also know the sampling probability of each record $p(r)$. We apply our data cleaning technique to the sample and get S_{clean} , and then apply our featurization to get features and labels (X_{clean}, Y_{clean}) . In the model update problem, we update the dirty model $\theta^{(d)}$ with some function $f(\cdot)$, i.e.,

$$\theta^{new} \leftarrow f(X_{clean}, Y_{clean}, \theta^{(d)})$$

Our goal is that these updates should minimize the error of the updated model and the true model $\theta^{(c)}$ (if we cleaned the entire data and trained over the entire data):

$$error(\theta^{new}) = \|\theta^{new} - \theta^{(c)}\|$$

5.2 Geometric Interpretation

We will present the update algorithm intuitively by describing it in terms of the convex geometry of the problem. Let us consider this problem in one dimension (i.e., the parameter θ is a scalar value), so then the goal is to find the minimum point (θ) of a curve $l(\theta)$. The consequence of dirty data is that we optimize the wrong loss function. We are optimizing the function $l'(\theta)$, when we really want to optimize $l(\theta)$. Figure 4A, illustrates the consequence of this optimization. The brown dotted line shows the loss function on the dirty data. If we optimize this loss function, we find the θ that at the minimum point. However, the true loss function (w.r.t to the clean data) is in blue. This optimal value is a suboptimal point on clean curve.




Figure 4: (A) A model trained on dirty data can be thought of as a sub-optimal point w.r.t to the clean data. (B) The gradient gives us the direction in which we need to move the suboptimal model to approach the true optimum.

We want to get to the optimal clean model $\theta^{(c)}$ which is visualized as a yellow star. The first question is which direction do we move θ . We do not know whether we need to move left or right. For this class of models, given a suboptimal point, we can find the direction to the global optimum. Mathematically, this direction corresponds to gradient of the loss function. The gradient is a function of the clean data that returns a d -dimensional vector w.r.t the current dirty optimal model i.e., the direction between the dirty model

and the true clean model. We need to move some distance γ along this direction (Figure 4b):

$$\theta^{new} \leftarrow \theta^{(d)} - \gamma \cdot \nabla \phi(\theta^{(d)})$$

In our visualization, this corresponds to whether we should move left or right. At the optimal point, the expected gradient will be zero. So intuitively, this approach iteratively moves the model downhill to, or corrects, the dirty model until the budget is reached.

However, in reality, we do not have the all of the clean data and have to approximate this gradient from a sample. The intuition, which we will formalize in Section 5.5, is that if we are on average in the right direction the algorithm is guaranteed to converge with analytics bounds on the convergence rate. It turns out that for an appropriately chosen γ , this convergence will hold for any batch size that we choose.

5.3 Average Gradient From a Sample

To derive a sample-based update rule, the first property that we should recognize is that sums commute with derivatives and gradients. Our class of Machine Learning models are based on loss minimization, that is they are a sum of losses, so the gradient $\nabla \phi(\theta)$ is:

$$\nabla \phi(\theta) = \frac{1}{n} \nabla \phi(x_i^{clean}, y_i^{clean}, \theta)$$

We estimate this gradient from a sample by averaging over a sample of clean data and reweighting them by their sampling probability. Let S be a sample of data, where each $i \in S$ is drawn with probability p_i :

$$\nabla \phi(\theta) \approx g_S(\theta) = \frac{1}{n |S|} \sum_{i \in S} \frac{1}{p_i} \nabla \phi(x_i^{clean}, y_i^{clean}, \theta)$$

Then for every batch of data cleaned, we apply the update to the current best model estimate:

$$\theta^{(t)} \leftarrow \theta^{(t-1)} - \gamma \cdot g_S(\theta^{(t-1)})$$

The detection in the previous step adds a small complication, since the sample S_{dirty} is not representative of all of the data. We also have to compensate for this bias by averaging this estimate with the gradient with respect to the clean data:

$$g_C(\theta) = \frac{1}{|R - R_{dirty}|} \sum_{i \in R - R_{dirty}} \nabla \phi(x_i^{clean}, y_i^{clean}, \theta)$$

Then, for weights α, β which we will subsequently discuss how to select (Section 5.4.1):

$$g(\theta) = \alpha \cdot g_C(\theta) + \beta \cdot g_S(\theta)$$

This average represents a novel optimization for the data cleaning that is not studied in the stochastic optimization literature. Since some data are less expensive to process than others, we can batch together processing the clean data and sample to process the dirty data. Finally, we add in the iteration, and at each iteration i , the update becomes:

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \gamma \cdot g(\theta^{(i)}) \blacksquare$$

5.4 Model Update Algorithm

We propose the iterative model correction algorithm. We initialize the algorithm with $\theta^{(d)}$ which is the dirty model. At each iteration $i = \{1, \dots, t\}$, we clean a batch of data b selected from the set of candidate dirty records R_{dirty} . Then, starting with $\theta^{(d)}$, we apply an averaged gradient update to get $\theta^{(i)}$. We iterate until our budget of cleaning $k = t \cdot b$ record is reached.

We present the model update algorithm here:

1. Calculate the gradient over the sample of clean data and call the result $g_S(\theta^{(i-1)})$
2. Calculate the average gradient over all the data in $R_{clean} = R - R_{dirty}$, and call the result $g_C(\theta^{(i)})$
3. Apply the following update rule:

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \lambda \cdot (\alpha \cdot g_S(\theta^{(i)}) + \beta \cdot g_C(\theta^{(i)}))$$

5.4.1 Selecting the Parameters

We have to set the parameters $b, \lambda, \alpha, \beta$. We will show that theory from the stochastic optimization literature can inform how to set these parameters. First, we will review all of the parameters that need to be set.

Batch Size b : The batch size b controls the frequency of iteration. Larger batches provide a more accurate estimate of the gradient at each iteration but analyst gets less frequent model updates.

Step Size γ : We have not explained how to pick the step size γ . In other words, how far should we travel in the gradient direction.

Weights α, β : The next problem is deriving the proportions with which we should combine $g_S(\theta)$ and $g_C(\theta)$. It turns out that $\alpha = \frac{R_{clean}}{R}$ and $\beta = \frac{R_{dirty}}{R}$, and we will show a derivation below.

Choosing S : As we mentioned, the optimal clean model depends on *all* the clean data, not just a sample. So g_S is an approximation of the true gradient g^* with some error $g^* \pm \epsilon$. The quality of the update depends on how well we can approximate g^* using g_S . The problem is how should we construct the sample of data to clean S to get the most accurate update. In particular, how should we choose the sampling probabilities p_i for each $i \in S$ such that the error is minimized.

5.5 Analysis with Stochastic Gradient Descent

This update policy can be formalized as a class of very well studied algorithms called Stochastic Gradient Descent. This gives us a theoretical framework to understand and analyze our update rule, bound the error, and choose points to clean. Mini-batch stochastic gradient descent (SGD) is an algorithm for finding the optimal value of θ , given the convex loss, and data. In mini-batch SGD, random subsets of data are selected at each iteration and the average gradient is computed for every batch. The key difference is that in traditional SGD there is no notion of dirty and clean data, and we have to formalize our method as a variant of mini-batch SGD to apply the theory.

ActiveClean as Lazy SGD: The update technique is ActiveClean can be formalized as a variant of SGD that lazily materializes the clean value. We apply SGD to dirty dataset, as data is sampled at each iteration, data is cleaned. To the SGD algorithm it is as if the entire data is clean but we are lazily materializing the clean values. It is well known that even for arbitrary initializations SGD makes significant progress in less than one epoch (a pass through the entire dataset) [8]. Furthermore in our setting, the dirty model can be much more accurate than an arbitrary initialization. This is the property that we exploit to make significant progress by cleaning only a sample of data.

Deriving α and β : We first want to show how to select α and β such that our estimate of the gradient is unbiased. We will then show that SGD will converge when the estimate is unbiased and the step-size is chosen appropriately. We first select a set of records before starting we select a set of dirty records $R_{dirty} \subseteq R$. We construct batch S_{dirty} only from the dirty records and apply cleaning to this batch. Now the problem is that this sample (and the resulting gradient) is possibly biased since we are excluding some data. However, since we know that the set $R_{clean} = R - R_{dirty}$ is clean, we can compute the average gradient over those records without cleaning:

$$g_C(\theta^t) = \frac{1}{|R - R_{dirty}|} \sum g_i(\theta^t)$$

Since we know the partition sizes, we can combine the two estimates g_C and g_S together:

$$g(\theta^t) = \frac{|R_{dirty}| \cdot g_S + |R_{clean}| \cdot g_C}{|R|}$$

Therefore,

$$\alpha = \frac{R_{\text{clean}}}{R}, \beta = \frac{R_{\text{dirty}}}{R}$$

LEMMA 1. *The gradient estimate $g(\theta^t)$ is unbiased if g_S is an unbiased estimate of:*

$$\frac{1}{|R_{\text{dirty}}|} \sum g_i(\theta^t)$$

PROOF SKETCH. This result follows directly from the linearity of expectation, since we are adding a deterministic result with an unbiased result. \square

Setting γ : There is extensive literature in machine learning for choosing the step size γ appropriately. γ can be set either to be a constant or decayed over time. Many machine learning frameworks (e.g., MLLib, Sci-kit Learn, Vowpal Wabbit) automatically set learning rates or provide different learning scheduling frameworks. In our experiments, we use a technique called inverse scaling where there is a parameter $\gamma_0 = 0.1$, and at each iteration we reduce it to $\gamma_t = \frac{\gamma_0}{|S|t}$.

Convergence: Convergence properties of batch SGD formulations has been well studied [11]. From the preceding analysis, we know that our gradient estimate is unbiased and we know how to select the step size.

PROPOSITION 1. *For an appropriately chosen learning rate γ_t , batch stochastic gradient descent will converge if $\mathbb{E}(g_S) = g^*$.*

Convergence Rate: The convergence rates of SGD are also well analyzed [7, 11, 46]. Suppose, a user cleans a batch size of b examples at each iteration. This allows us to bound the error of intermediate models and understand the expected number of steps before a model within a certain error.

PROPOSITION 2. *For a general convex loss, a batch size b , and t iterations, the convergence rate is bounded by $O(\frac{\sigma^2}{\sqrt{bt}})$. σ^2 is the variance in the estimate of the gradient at each iteration:*

$$\mathbb{E}(\|g_S - g^*\|^2)$$

Setting the batch size: The batch size should be set by the analyst to have the desired properties. Larger batches will take longer to clean and will make more progress towards the clean model, but will have less frequent model updates. On the other hand, smaller batches are cleaned faster and have more frequent model updates but make less progress. The overheads introduced by our approach are more evident at smaller batch sizes. There are diminishing returns to increasing the batch size $O(\frac{1}{\sqrt{b}})$. Empirically, we find that batch sizes of 50 converge the fastest on our data and we specify this parameter in our experiments. If a data cleaning technique requires a larger batch size than this, i.e., data cleaning is fast enough that the iteration overhead is significant compared to cleaning 50 records, we can always apply the updates in smaller batches. For example, the batch size set by the analyst might be $b = 1000$, but we apply the model updates after every 50 records are cleaned. This way we can dissociate the requirements of SGD and the data cleaning technique.

Non-convex losses: We acknowledge that there is an increasing popularity of non-convex losses in the Neural Network and Deep Learning literature. However, even for these losses, gradient descent techniques still apply. Instead of converging to a global optimum they converge to a locally optimal value. Likewise, ActiveClean will converge to the closest locally optimal value to the dirty

model. Because of this, it is harder to reason about the results. Different initializations will lead to different local optima, and thus, introduces a complex dependence on the initialization with the dirty model. This problem is not fundamental to ActiveClean and any gradient technique suffers this challenge for general non-convex losses, and we hope to explore this more in the future.

6. SAMPLING

In the previous section, we assumed that our model update received a sample with probabilities $p(r)$. In this section, we derive an optimal sampling problem that directly follows from our analysis of our update rule via SGD. It will turn out that the solution to the optimal sampling problem is not realizable in practice (as it depends on knowing the cleaned value), but we can use this to inform the next section where we estimate the cleaned value.

6.1 Goals and Challenges

In the Machine Learning and Optimization literature, SGD algorithms are optimized to avoid scanning the entire data. Uniform sampling is cheap so it is the preferred solution. However, data cleaning costs can be many orders of magnitude higher than model training. As a result, uniform sampling may not be the most efficient option. We can sacrifice computational overhead by precomputing some results over the entire data for savings during the data cleaning phase. We formulate this problem as an optimal sampling problem where we want to compute the sampling probabilities $p(r)$ that maximize the accuracy of our updates.

6.2 Optimal Sampling Problem

Recall that the convergence rate of an SGD algorithm is bounded by σ^2 which is the variance of the gradient. Intuitively, the variance measures how accurately we estimate the gradient from a uniform sample. Other sampling distributions, while preserving the sample expected value, may have a lower variance. Thus, we define the optimal sampling problem as a search over sampling distributions to find the minimum variance sampling distribution.

DEFINITION 3 (OPTIMAL SAMPLING PROBLEM). *Given a set of candidate dirty data R_{dirty} , $\forall r \in R_{\text{dirty}}$ find sampling probabilities $p(r)$ such that over all samples S of size k it minimizes:*

$$\mathbb{E}(\|g_S - g^*\|^2)$$

In other words, we want to most accurately (in the mean squared error sense) estimate the gradient.

To construct these sampling probabilities, we first need the following lemma about importance sampling. This lemma describes the optimal distribution over a set of scalars:

LEMMA 2. *Given a set of real numbers $A = \{a_1, \dots, a_n\}$, let \hat{A} be a sample with replacement of A of size k . If μ is the mean \hat{A} , the sampling distribution that minimizes the variance of μ , i.e., the expected square error, is $p(a_i) \propto a_i$.*

PROOF SKETCH. This proof follows from [30], as it is a straight forward importance sampling result. It is easy to verify that for sampling probabilities $p(a_i)$, that the unbiased estimate of the mean is:

$$\mu = \frac{1}{nk} \cdot \sum_{i \in \hat{A}} \frac{a_i}{p_i}$$

If we minimize the variance $\text{Var}(\mu)$, we can find that:

$$p_i = \frac{|a_i|}{\sum_i |a_i|}$$

\square

Lemma 2 shows that when estimating a mean of numbers with sampling, the distribution with optimal variance is where we sample proportionally to the values. This insight leads to a direct higher-dimensional generalization, where at iteration t we should sample the records in R_{dirty} with probabilities:

$$p_i \propto \|\nabla\phi(x_i^{(clean)}, y_i^{(clean)}, \theta^t)\|$$

However, in our case, it leads to a chicken-and-egg problem. The optimal sampling distribution requires knowing $(x_i^{(clean)}, y_i^{(clean)})$, however, we have to sample and clean those points to get those values. In the next section, we discuss how to inexpensively approximate this optimal distribution. As our technique can work with *any* distribution, we are guaranteed convergence no matter how inaccurate this approximation is. However, a better approximation will lead to an improved convergence rate.

7. ESTIMATION

In this section, we make the sampling result of the previous section practical, and estimate the cleaned data.

7.1 Challenges and Goal

The optimal sampling distribution is dependent on a value that we cannot know without data cleaning $\nabla\phi(x_i^{(clean)}, y_i^{(clean)}, \theta^t)$. One way to approximate this distribution is to learn a function $e(\cdot)$ mapping (x_{clean}, y_{clean}) , based on the data that we have cleaned. Suppose, we use a regression technique based on the previous data. This is a high-dimensional regression problem which may have to learn a very complicated relationship between dirty and clean data. The biggest challenge with such an estimator is the cold start problem, where if we have cleaned a small amount of data, the estimator will be inaccurate. In ActiveClean, we want to be able to make as much progress as possible in the early iterations so this technique may not work. We take an alternative approach, where we try to exploit what we know about data cleaning to produce an estimate for groups of similarly corrupted records. To define ‘‘similarly corrupted’’, we are going to show how we can use the detection step to make this problem tractable.

7.2 Estimation For A Priori Detection

EXAMPLE 6. Consider our EEG dataset with multiple types of errors. Suppose the electrical signals for some records are corrupted with missing data and the patient data for other records are corrupted with inconsistency problems. Each training example will have a set of corrupted features (e.g., $\{1, 2, 6\}$, $\{1, 2, 15\}$).

Suppose that we have just cleaned the records r_1 and r_2 represented as tuples with their corrupted feature set: $(r_1, \{1, 2, 3\})$, $(r_2, \{1, 2, 6\})$. Then, we have a new record $(r_3, \{1, 2, 3, 6\})$. We want to be able to use the cleaning results from r_1, r_2 to estimate the gradient in r_3 .

If most of the features are correct, it would seem like the gradient is only incorrect in one or two of its components. The problem is that the gradient $\nabla\phi(\cdot)$ can be a very non-linear function of the features that couple features together. For example, let us look at the gradient for linear regression:

$$\nabla\phi(x, y, \theta) = (\theta^T x - y)x$$

We see that it is not possible to isolate the effect of a change of one feature on the gradient. Even if one of the features is corrupted, all of the gradient components will be incorrect.

7.2.1 Error Decoupling

We can try to approximate the gradient in a way that the effects of features on the gradient are decoupled. Recall, that when we formalized the a priori detection problem, we ensured that associated with each $r \in R_{dirty}$ is a set of errors f_r which is a set that

identifies a set of corrupted features. We will show how we can use this property to construct a coarse estimate of the clean value. The main idea is that if we can calculate average changes for each feature, then given an uncleaned (but dirty) record, we can add these average changes to correct the gradient.

Let us formalize this intuition. Instead of computing the actual gradient with respect to the true clean values, let us compute the conditional expectation given that a set of features and labels f_r are corrupted:

$$p_i \propto \mathbb{E}(\nabla\phi(x_i^{(clean)}, y_i^{(clean)}, \theta^t) \mid e_r)$$

What we mean by corrupted features is that:

$$i \notin f_r \implies x_{clean}[i] - x_{dirty}[i] = 0$$

$$i \notin f_r \implies y_{clean}[i] - y_{dirty}[i] = 0$$

The needed approximation represents a linearization of the errors. We will show that the sampling distribution can be estimated in the form:

$$p_r \propto \|\nabla\phi(x, y, \theta^t) + M_x \cdot \Delta_{rx} + M_y \cdot \Delta_{ry}\|$$

where M_x, M_y are matrices and Δ_{rx} and Δ_{ry} are average change vectors for the corrupted features in r . Without decoupling, if we were to calculate the expected value conditioned on e_r , we would have to condition on all the combinatorial possibilities.

7.2.2 Deriving M_x, M_y

We can take the expected value of the Taylor series expansion around the dirty value. If d is the dirty value and c is the clean value, the Taylor series approximation for a function f is given as follows:

$$f(c) = f(d) + f'(d) \cdot (d - c) + \dots$$

If we ignore the higher order terms, we see that the linear term $f'(d) \cdot (c - d)$ decouples the features. We only have to know the change in each feature to estimate the change in value. In our case the function f is the gradient $\nabla\phi$. So, the resulting linearization is:

$$\begin{aligned} \nabla\phi(x_i^{(clean)}, y_i^{(clean)}, \theta) &\approx \nabla\phi(x, y, \theta) + \frac{\partial}{\partial X} \nabla\phi(x, y, \theta) \cdot (x - x_{clean}) \\ &\quad + \frac{\partial}{\partial Y} \nabla\phi(x, y, \theta) \cdot (y - y_{clean}) \end{aligned}$$

When we take the expected value:

$$\begin{aligned} \mathbb{E}(\nabla\phi(x_{clean}, y_{clean}, \theta)) &\approx \nabla\phi(x, y, \theta) + \frac{\partial}{\partial X} \nabla\phi(x, y, \theta) \cdot \mathbb{E}(\Delta x) \\ &\quad + \frac{\partial}{\partial Y} \nabla\phi(x, y, \theta) \cdot \mathbb{E}(\Delta y) \end{aligned}$$

So the resulting estimation formula takes the following form:

$$\approx \nabla\phi(x, y, \theta) + M_x \cdot \mathbb{E}(\Delta x) + M_y \cdot \mathbb{E}(\Delta y)$$

Recall that we have a d dimensional feature space and l dimensional label space. Then, $M_x = \frac{\partial}{\partial X} \nabla\phi$ is a $d \times d$ matrix, and $M_y = \frac{\partial}{\partial Y} \nabla\phi$ is a $d \times l$ matrix. Both of these matrices are computed with respect to dirty data, and we will present an example. Δx is a d dimensional vector where each component represents a change in that feature and Δy is an l dimensional vector that represents the change in each of the labels.

7.2.3 Example M_x, M_y

Let us walk through an example, where we calculate M_x, M_y . In the linear regression example, where the gradient is:

$$\nabla\phi(x, y, \theta) = (\theta^T x - y)x$$

For a record, r , suppose we have a feature vector x . If we take the partial derivatives with respect to x, M_x is:

$$M_x[i, i] = 2x[i] + \sum_{i \neq j} \theta[j]x[j] - y$$

$$M_x[i, j] = \theta[j]x[j]$$

Similarly M_y is:

$$M_y[i, 1] = x[i]$$

In the appendix, we describe the matrices for common convex losses.

7.2.4 More Accurate Early Error Estimates

We can give some intuition on why the linearization actually avoids amplifying estimation error. If we consider the linear regression gradient:

$$\nabla\phi(x, y, \theta) = (\theta^T x - y)x$$

We can rewrite this as a vector in each component:

$$g[i] = \sum_i x[i]^2 - x[i]y + \sum_{j \neq i} \theta[j]x[j]$$

We see that this function is already mostly linear in x except for the one quadratic term. However, this one quadratic term has potential to amplify errors. Consider two expressions:

$$\begin{aligned} f(x + \epsilon) &= (x + \epsilon)^2 = x^2 + 2x\epsilon + \epsilon^2 \\ f(x + \epsilon) &\approx f(x) + f'(x)(\epsilon) = x^2 + 2x\epsilon \end{aligned}$$

The only difference between the two estimates is the quadratic ϵ^2 , if ϵ is highly uncertain random variable then the quadratic dominates. Our estimate ϵ is an average, so at the least, its error is given by the variance in the data. If this variance is large, the Taylor estimate avoids amplifying this error. We evaluate our technique in Section 8.5 against alternatives, and we find that indeed we provide more accurate estimates for a small number of samples cleaned. When the number of cleaned samples increases the alternative techniques are comparable.

7.2.5 Maintaining Decoupled Averages

This linearization allows us to maintain per feature (or label) average changes and use these changes to center the optimal sampling distribution around the expected clean value. We know how to estimate $\mathbb{E}(\Delta x)$ and $\mathbb{E}(\Delta y)$.

LEMMA 3 (SINGLE FEATURE). *For a feature i , we average all records cleaned that have an error for that feature, weighted by their sampling probability:*

$$\bar{\Delta}_i = \frac{1}{NK} \sum_{j=1}^K (x[i] - x_{clean}[i]) \times \frac{1}{p_j}$$

Similarly, for a label i :

$$\bar{\Delta}_i = \frac{1}{NK} \sum_{j=1}^K (y[i] - y_{clean}[i]) \times \frac{1}{p_j}$$

Then, it follows, that we can aggregate the $\bar{\Delta}_i$ into a single vector:

LEMMA 4 (DELTA VECTOR). *Let $\{1.., i, ..., d\}$ index the set of features and labels. For a record r , the set of corrupted features is f_r . Then, each record r has a d -dimensional vector Δ_r which is constructed as follows:*

$$\Delta_r[i] = \begin{cases} 0 & i \notin f_r \\ \bar{\Delta}_i & i \in f_r \end{cases}$$

With the above theorem, we finally have an approximation to our sampling weights:

$$p_r \propto \|\nabla\phi(x, y, \theta^{(t)}) + M_x \cdot \Delta_{rx} + M_y \cdot \Delta_{ry}\| \blacksquare$$

7.3 Estimation For Adaptive Case

The same logic still holds in the adaptive setting, however, we have to reformulate what we mean by ‘‘similarly’’ corrupted. Here, we use u dirtiness classes. Instead of conditioning on the features that are corrupted, we condition the classes. So for each error class, we compute a Δ_{ux} and Δ_{uy} . These are the average change in the features given that class and the average change in labels given that class respectively.

$$p_{r,u} \propto \|\nabla\phi(x, y, \theta^{(t)}) + M_x \cdot \Delta_{ux} + M_y \cdot \Delta_{uy}\| \blacksquare$$

8. EXPERIMENTS

There are a number of different axes on which we can evaluate ActiveClean. First, we take real datasets and generate various types of errors to illustrate the value of data cleaning in comparison to robust statistical techniques. Next, we explore different prioritization and model update schemes for data cleaning samples. Finally, we evaluate ActiveClean end-to-end in a number of real-world data cleaning scenarios.

8.1 Experimental Setup and Notation

Our main metric for evaluation is a relative measure of trained model with ActiveClean (or an alternative technique) and the model if all of the data is cleaned.

Relative Model Error. Let θ be the model trained on the dirty data, and let θ^* be the model trained on the same data if it was cleaned. Then the model error is defined as $\frac{\|\theta - \theta^*\|}{\|\theta^*\|}$.

8.1.1 Scenarios

We apply ActiveClean to the following scenarios:

Income Classification (Adult): In this census dataset, our task is to predict the income bracket (binary) from 12 numerical and categorical covariates. There are 45552 data points in this dataset. We use a SVM classifier to predict the income bracket of the person.

Seizure Classification (EEG): In this dataset, our task is to predict the on set of a seizure (binary) from 15 numerical covariates. There are 14980 data points in this dataset. We chose this dataset since the classification task is hard with an accuracy in clean data of 65%. The model that we use is a thresholded Linear Regression.

Handwriting Recognition (MNIST): In this dataset, our task is to classify 60,000 images of handwritten images into 10 categories. The unique part of this dataset is the featurized data consists of a 784 dimensional vector which includes edge detectors and raw image patches. We use this dataset to explore how we can corrupt the raw data to affect subsequent featurization. The model is an one-to-all multiclass SVM classifier.

Dollars For Docs: ProPublica has collected a dataset of 240089 corporate contributions to doctors for analysis of whether these contributions affect research. There are 5 textual attributes and one numerical attribute. We use this dataset to classify whether a research contribution is ‘‘not covered’’ under the declared experimental protocol from the attributes. The model is a binary SVM using a bag-of-words model to featurize the textual attributes.

8.1.2 Compared Algorithms

Here are the alternative methodologies that we consider:

Robust Logistic Regression [16]. Feng et al. proposed a variant of logistic regression that is robust to outliers. We chose this algorithm because it is a robust extension of the convex regularized loss model, leading to a better apples-to-apples comparison between the techniques.

Discarding Dirty Data. As a baseline we explore model accuracy when dirty data is discarded.

SampleClean (SC) [43]. In SampleClean, we take a sample of data, apply data cleaning, and then train a model to completion.

Active Learning (AL) [19]. There have been recent proposals of integrating Active Learning with stochastic optimization. Active Learning has been widely applied in the data cleaning literature [18], but never integrated with model training. We acknowledge the work in the Machine Learning literature and formulate an Active Learning that uses a prioritization function agnostic of the dirty data

(no estimation and detection). We clean in batches of 50 records and we use this parameter setting for both AL and ActiveClean.

ActiveClean Oracle (AC+O): In ActiveClean Oracle, we importance sample points by their clean gradient. This represents the theoretical best that our algorithm could hope to achieve given perfect estimation.

8.2 Does Data Cleaning Matter

Before we evaluate ActiveClean, we first evaluate the benefits of cleaning on 2 of our example datasets (EEG and Adult). We first explore this problem without sampling to understand which types of data corruption are amenable to data cleaning and which are better suited for robust statistical techniques. We compare 4 schemes: (1) full data cleaning uncorrupted, (2) baseline of no cleaning, (3) discarding the dirty data, and (4) robust logistic regression. We corrupted 5% of the training examples in each dataset in two different ways:

Random Corruption: We simulated high-magnitude random outliers. We select 5% of the examples and features uniformly at random and replace a feature with 3 times the highest feature value.

Systematic Corruption: We simulated innocuous looking (but still incorrect) systematic corruption. We trained the model on the clean data, find the three most important feature (highest weighted). We sort examples by each these features and corrupt the top of examples with the mean value for that feature. At the end, we have corrupted 5% of the examples. It is important to note the some examples can have multiple corrupted features.




Figure 5: (a) Robust techniques and discarding data work when corrupted data are random and look atypical. (b) Data cleaning can provide reliable performance in both the systematically corrupted setting and randomly corrupted setting.

We plot the test accuracy for models trained on both types of data with the different techniques, since it is a binary classification task with balanced class distributions we cut the plots at 50% (random accuracy) (Figure 5). As we argued in this paper, the robust method performs well on the random high-magnitude outliers, however, falters on the systematic corruption. In the random setting, discarding dirty data also performs well. However, when the corruption is systematic, neither alternative works. Data cleaning is the most reliable option across datasets and corruption types. The problem is that without cleaning we do not know if the corruption is random or systematic. While data cleaning requires more effort, it provides benefits in both settings. In the remaining experiments, unless otherwise noted, we use the systematic corruption.

Summary: Data cleaning mitigates the effects of both systematic and random corruption.

8.3 ActiveClean A Priori Detection

The next set of experiments evaluate different approaches to cleaning a sample of data. In this set of experiments, we use the *a priori* case for detection. We assume that we know all of the corrupted records in advance but do not know how to repair them.

8.3.1 Comparison with Active Learning and SampleClean

In our first experiment, we evaluate the samples-to-error trade-off between three alternative algorithms: ActiveClean (AC), SampleClean, Active Learning, and ActiveClean + Oracle (AC+O). In Figure 6, we present our results on Adult and EEG. We find that ActiveClean gives its largest benefits for small sample sizes (up-to 12x more accurate than SampleClean). ActiveClean makes significant progress because of its intelligent initialization, iterative updates, and partitioning. For example, the EEG dataset is the hardest classification task. SampleClean has difficulty on this dataset since it takes a uniform sample of data (only 5% of which are corrupted on average) and tries to train a model using only this data. ActiveClean and Active Learning leverage the initialization from the dirty data to get an improved result. However, ActiveClean and Active Learning differ in the way they prioritize cleaning. Active Learning prioritizes with respect to a dirty model (i.e., is agnostic to data error) and ActiveClean uses estimation and detection. ActiveClean’s estimates and detection allow us to beat Active Learning on all three of the datasets. ActiveClean also is relatively close in performance to the oracular version (theoretical optimum).

However, to an end user the metric that matters is test accuracy. In Figure 7, we present the results for the two datasets: Adult and EEG. We find that in both Adult and EEG, ActiveClean converges to clean test accuracy faster than the alternatives.

Summary: ActiveClean with a priori detection converges faster than SampleClean and Active Learning.




Figure 6: ActiveClean converges with a smaller sample size to the true result in comparison to Active Learning and SampleClean. We show the relative model error as a function of the number of examples cleaned.




Figure 7: ActiveClean converges with a smaller sample size to the maximum test accuracy in comparison to Active Learning and SampleClean. The reductions in model error correlate well with increased test accuracy in (a). Due to the hardness of the task in (b), we small gains in terms of test accuracy.

8.3.2 Source of Improvements

Now, we try to understand the source of our improvements w.r.t Active Learning and SampleClean. We pick a single point on the curves shown in Figure 6 that corresponds to 500 records cleaned

and compare the performance of ActiveClean with and without various optimizations. This is a vertical slice of the plots in the previous experiments.

We denote ActiveClean without detection as (AC-D) (that is at each iteration we sample from the entire dirty data) and ActiveClean without detection and importance sampling as (AC-D-I). In Figure 8, we plot the relative error of the alternatives and ActiveClean with and without the optimizations. Detection significantly improves our results in all of the datasets, and accounts for a substantial part of the improvements over Active Learning. However, when we remove detection we still see some improvements since our importance sampling relies on error impact estimates. Not surprisingly, when we remove both these optimizations, ActiveClean is slightly worse (but still comparable to) than Active Learning.

Summary: Both a priori detection and non-uniform sampling significantly contribute to the gains over Active Learning.




Figure 8: We clean 500 records and plot the relative error of ActiveClean with and without optimizations. -D denotes no detection, and -D-I denotes no detection and no importance sampling. Both optimization significantly help ActiveClean outperform SampleClean and Active Learning.

8.3.3 Corruption Rate

Both Active Learning and ActiveClean outperform SampleClean in our experiments. In our next experiment (Figure 9), we try to understand how much of this performance is due to the initialization (i.e., SampleClean trains a model from “scratch”). We vary the systematic corruption rate and plot the number of records cleaned to achieve 1% relative error for SampleClean and ActiveClean. SampleClean does not use the dirty data and thus is not dependent on this rate. We find that SampleClean outperforms ActiveClean only when corruptions are very severe (45% in Adult and nearly 60% in EEG).

Summary: SampleClean is beneficial in comparison to ActiveClean when corruption rates exceed 45% in our experiments.




Figure 9: We corrupt an increasing number of entries in the data matrix. ActiveClean performs well until the corruption is so severe that the dirty model is not a good initialization.

8.4 Adaptive Detection

In this experiment, we explore how the results of the previous experiment change when we use an adaptive detector rather than




Figure 10: Even with a classifier ActiveClean converges faster than Active Learning and SampleClean.

the a priori detector. Recall, in our systematic corruption, we corrupted 3 of the most informative features at random, thus we group these problems into 3 classes. We use an all-versus-one SVM to learn the categorization.

8.4.1 Basic Performance

In Figure 10, we overlay the convergence plots in the previous experiments with a curve (denoted by AC+C) that represents ActiveClean using a classifier instead of the a priori detection. We find an interesting tradeoff where initially ActiveClean is comparable to Active Learning, as our classifier becomes more effective the detection improves the performance. For both datasets, after cleaning 1000 records, ActiveClean is within 10% of a priori detection performance.

Summary: ActiveClean with a classifier converges faster than Active Learning and SampleClean.

8.4.2 Classifiable Errors

The adaptive case depends on being able to predict corrupted records. For example, random corruption that look like other data may be hard to learn. As corruption becomes more random, the classifier becomes increasingly erroneous. We run an experiment where we start with the systematic corruption described earlier. We increasingly make this corruption more random. Instead of selecting the highest valued records for the most valuable features, we corrupt random records with probability p . We compare these results to AC-D where we do not have a detector at all at one vertical slice of the previous plot (cleaning 1000 records). In Figure 11, we plot the relative error reduction using a classifier. We find that when the corruption is about 50% random then we reach a break even point.

Summary: The adaptive detection can tolerate some misclassifications and still provides improvements until corruption is very random.




Figure 11: Data corruptions that are less random are easier to classify, and lead to more significant reductions in relative model error.

8.5 Estimation

In Figure 12, we explore how estimation technique performs compared to a few alternative estimation techniques: (1) “linear regression” where we train a linear regression model that predicts

the clean gradient as a function of the dirty gradient, (2) “average gradient” change where we do not use the detection to inform how to apply the estimate, (3) “average feature change” where we do not linearize but use detection, and (4) our Taylor series linear approximation. We only how accurately the technique can estimate the gradient for sampling. More accurate gradient estimates will result in faster convergence.

Using the EEG dataset, we measure the relative estimation error (relative L2 error with the true gradient) for different amounts of data cleaned. The Taylor series approximation proposed in this work gives more accurate for small cleaning sizes. We give the intuition in Section 7.2.4 where the linearization prevents amplification of errors due to quadratic terms. Linear regression and the average feature change technique do eventually perform comparably but only after cleaning much more data.

Summary: Linearized gradient estimates are more accurate when estimated from small samples.




Figure 12: The Taylor series approximation gives more accurate estimates when the amount of cleaned data is small.

8.6 Real Scenarios

We evaluate ActiveClean in two real scenarios, one demonstrating the a priori case and one for the adaptive detection case.

8.6.1 A Priori: Constraint Cleaning

In our first real scenario, we explore the Dollars for Docs dataset published by Pro Publica. This dataset consists of medical research contributions from pharmaceutical corporations and subsequent research by universities and hospitals. In this dataset, our application is inspired by a fraud detection application. Each research publication is with whether the contribution is “covered” or “not covered” by the experimental protocol. Our goal is to classify research contributions as being not covered (i.e., possibly fraudulent) based on 5 attributes: the principal investigator’s speciality, the drug or device in question, the corporation, the amount contributed, and whether the research was disputed.

However, the dataset is very dirty, and Pro Publica highlights the cleaning that they performed before publishing this dataset [2] including entity resolution and enforcing quality constraints. Errors include:

Corporations are inconsistently represented: “Pfizer”, “Pfizer Inc.”, “Pfizer Incorporated”.

Drugs are inconsistently represented: “TAXOTERE DOCETAXEL-PROSTATE CANCER” and “TAXOTERE”

Label of covered and not covered are not consistent: “No”, “Yes”, “N”, “This study is not supported”, “None”, “Combination”

Research subject must be a drug OR a medical device and not both: “BIO FLU QPAN H7N9AS03 Vaccine” and “BIO FLU QPAN H7N9AS03 Device”

We treat this problem as a constraint-based cleaning problem where we encode these problems as data quality constraints.

Rule 1: Matching dependency on corporation (Weighted Jaccard Similarity > 0.8).

Rule 2: Matching dependency on drug (Weighted Jaccard Similarity > 0.8).

Rule 3: Label must either be “covered” or “not covered”.

Rule 4: Either drug or medical device should be null.

This provides us with a detector for the corruption. We confirmed drug names with company catalogs and used judgement to make the labels consistent. To run this experiment, we manually cleaned the entire dataset up front, and simulated sampling from the dirty data and cleaning by looking up the value in the cleaned data. Figure 13a shows that ActiveClean converges faster than Active Learning and SampleClean. To achieve a 4% relative error (i.e., a 75% error reduction from the dirty model), ActiveClean cleans 40000 fewer records than Active Learning. Also, for 10000 records cleaned, ActiveClean has 3.5x smaller error than SampleClean.

Figure 13b shows the true positive rate (fraction of non-covered research contributions identified) of the classifier as a function of the number of records cleaned. We find that on the dirty data, we can only classify 66% of the positive examples correctly. On the cleaned data, this classifier is nearly perfect with a 97% true positive rate. ActiveClean converges to the cleaned accuracy faster than the alternatives.

Summary: In a real a priori detection scenario, ActiveClean significantly reduces the number of records to clean to train an accurate model.




Figure 13: (A) The relative model error as a function of the number of cleaned records. (B) The true positive rate as a function of the number of cleaned records.

8.6.2 Adaptive: Replacing Corrupted Handwriting Data

We consider the following scenario with the MNIST handwritten digit recognition dataset. Suppose, our goal is to classify digits into a set of classes. However, we suspect that some of our raw images are of low quality. Typically image processing workflows are long with many different featurization steps. As a result, changes to the raw images may have very significant effects on some features. In this scenario, the analyst must inspect a potentially corrupted image and replace it with a higher quality one. We devise an “adaptive detection” ActiveClean scenario for this example.

The MNIST dataset consists of 64x64 grayscale images. We run two experiments, in which we have two types of corruptions: (1) 5x5 block removal where take a random 5x5 block from the image and set its pixel values to 0, and (2) Fuzzy where we run a 4x4 moving average over the entire image. We applied these corruptions to a random 5% of the images. We constructed these corruptions to mimic the random vs. systematic corruption that we studied before. The 5x5 block removal behaves much more like a systematic corruption. Typical image processing features are based on edges and corrupting edges leads to ambiguities. On the other hand, the making the image fuzzy is more like a random corruption. We use a 9 class classifier (one for each digit) to detect the corruption.

Figure 14 shows that ActiveClean makes more progress towards the clean model with a smaller number of examples cleaned. To achieve a 2% error for the block removal, we can inspect 2200

fewer images than Active Learning and 2750 fewer images than SampleClean. For the fuzzy images, both Active Learning and ActiveClean reach 2% error after cleaning fewer than 100 images, while SampleClean requires 1750.

Summary: In a real adaptive detection scenario, ActiveClean significantly reduces the number of records to clean to train an accurate model.




Figure 14: In a real adaptive detection scenario with the MNIST dataset, ActiveClean outperforms Active Learning and SampleClean. We simulate systematic corruptions with image block removal and random corruptions with image fuzzifying.

9. RELATED WORK

Bringing together data cleaning and machine learning presents us with several exciting new research opportunities that incorporates results from both communities. We highlight some of the key relevant work in this field and how this relates to our proposal.

Stochastic Optimization and Active Learning: Zhao and Tong recently proposed using importance sampling in conjunction with stochastic gradient descent [46]. However, calculating the optimal importance sampling distribution is very expensive and is only justified in our case because data cleaning is even more expensive. Zhao and Tong use a different approximation to work around this problem. This work is one of many in an emerging consensus in stochastic optimization that not all data are equal (e.g., [33]). This line of work builds on prior results in linear algebra that show that some matrix columns are more informative than others [13], and Active Learning which shows that some labels are more informative than others [36]. Active Learning largely studies the problem of label acquisition [36], and recently the links between Active Learning and Stochastic optimization have been studied [19]. We use the work in Guillory et al. to evaluate a state-of-the-art Active Learning technique against our method.

Transfer Learning and Bias Mitigation: ActiveClean has a strong link to a field called Transfer Learning [31]. The basic idea of Transfer Learning is that suppose a model is trained on a dataset D but tested on a dataset D' . In transfer learning, the model is often weighted or transformed in such a way that it can still predict on D' . Transfer Learning has not considered the data cleaning setting, in which there is a bijective map between $D \mapsto D'$ that is expensive to compute. Much of the complexity and contribution of our work comes from efficiently cleaning the data. In robotics, Mahler et al. explored a calibration problem in which data was systematically corrupted [27] and proposed a rule-based technique for cleaning data. Other problems in bias mitigation (e.g., Krishnan et al. [26]) have the same structure, systematically corrupted data that is feeding into a model. In this work, we try to generalize these principles given a general dirty dataset, convex model, and data cleaning procedure.

Secure Learning: Another relevant line of work is the work in private machine learning [14,41]. Learning is performed on a noisy variant of the data which mitigates privacy concerns. The goal is to extrapolate the insights from the noisy data to the hidden real data. Our results are applicable in this setting in the following way.

Imagine, we were allowed to query k true data points from the real data, which points are the most valuable to query. This is also related work in adversarial learning [29], where the goal is to make models robust to adversarial data manipulation.

Data Cleaning and Databases: There are also several recent results in data cleaning that we would like to highlight. Progressive data cleaning methodologies have been proposed, however, these techniques tend to be application agnostic [28]. Altowim et al. proposed a framework for progressive entity resolution [5]. Volkovs et al. explored a related topic of maintaining data cleaning rules that change over time [40]. Recently, in works such as SampleClean [43], the application (i.e., queries) are used to inform data cleaning methodology. When the workload is made up of aggregate queries, cleaning samples of data may suffice. Similarly, Bergman et al. explore the problem of query-oriented data cleaning [6]. Given a query they clean data relevant to that query. Bergman et al. does not explore the Machine Learning applications studied in this work. Deshpande et al. studied data acquisition in sensor networks [12]. They explored value of information based prioritization of data acquisition for estimating aggregate queries of sensor readings. Finally, incremental optimization methods like SGD have a connection to incremental materialized view maintenance as the argument for incremental maintenance over recomputation is similar (i.e., relatively sparse updates). Krishnan et al. explored how samples of materialized views can be maintained similar to how models are updated with a sample of clean data in this work [35]

10. CONCLUSION

As analytics frameworks increasingly support Machine Learning, methodologies for optimizing across the entire pipeline from data cleaning to model training are required. In this paper, we propose ActiveClean, a framework that integrates existing data cleaning solutions with Machine Learning model training. ActiveClean information from a statistical model to inform data cleaning on samples, and uses the information from clean samples to update a statistical model. The tight feedback loop between model training and data cleaning poses several new algorithmic challenges which we address with a stochastic gradient descent framework. Our experimental results are really promising as they suggest that these optimization can significantly reduce data cleaning costs. Techniques such as Active Learning and SampleClean are not optimized for this setting, and we are able to achieve sharper cost-accuracy tradeoffs than either technique.

ActiveClean is only a first step in a larger integration of data analytics and data acquisition/cleaning. There are several exciting, new avenues for future work. The empirical success of Deep Learning has led to increasing industry and research adoption in many tasks that were traditionally served by convex models. We hope to explore how we can integrate with such frameworks and what we can analyze about the performance. We will also explore what we can do if the analytics at the end of the pipeline is a black box function. Techniques proposed in Reinforcement Learning likely apply as we simultaneously have to learn the properties of the function as we are cleaning.

11. REFERENCES

- [1] Big data's dirty problem. <http://fortune.com/2014/06/30/big-data-dirty-problem/>.
- [2] Dollars for docs. <http://projects.propublica.org/open-payments/>.
- [3] List of countries by imports. https://en.wikipedia.org/wiki/List_of_countries_by_imports.
- [4] Sampleclean. <http://sampleclean.org/>, 2015.
- [5] Y. Altowim, D. V. Kalashnikov, and S. Mehrotra. Progressive approach to relational entity resolution. *Proceedings of the VLDB Endowment*, 7(11), 2014.

- [6] M. Bergman, T. Milo, S. Novgorodov, and W.-C. Tan. Query-oriented data cleaning with oracles. 2015.
- [7] D. P. Bertsekas. Incremental gradient, subgradient, and proximal methods for convex optimization: A survey. *Optimization for Machine Learning*, 2010:1–38.
- [8] L. Bottou. Stochastic gradient descent tricks. In *Neural Networks: Tricks of the Trade*, pages 421–436. Springer, 2012.
- [9] X. Chu, J. Morcos, I. F. Ilyas, M. Ouzzani, P. Papotti, N. Tang, and Y. Ye. KATARA: A data cleaning system powered by knowledge bases and crowdsourcing. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data, Melbourne, Victoria, Australia, May 31 - June 4, 2015*, pages 1247–1261, 2015.
- [10] A. Crotty, A. Galakatos, and T. Kraska. TUPLEWARE: Distributed machine learning on small clusters. *IEEE Data Eng. Bull.*, 37(3):63–76, 2014.
- [11] O. Dekel, R. Gilad-Bachrach, O. Shamir, and L. Xiao. Optimal distributed online prediction using mini-batches. *The Journal of Machine Learning Research*, 13(1):165–202, 2012.
- [12] A. Deshpande, C. Guestrin, S. Madden, J. M. Hellerstein, and W. Hong. Model-driven data acquisition in sensor networks. In *(e)Proceedings of the Thirtieth International Conference on Very Large Data Bases, Toronto, Canada, August 31 - September 3 2004*, pages 588–599, 2004.
- [13] P. Drineas, M. Magdon-Ismail, M. W. Mahoney, and D. P. Woodruff. Fast approximation of matrix coherence and statistical leverage. *The Journal of Machine Learning Research*, 13(1):3475–3506, 2012.
- [14] J. C. Duchi, M. I. Jordan, and M. J. Wainwright. Local privacy and statistical minimax rates. In *Foundations of Computer Science (FOCS), 2013 IEEE 54th Annual Symposium on*, pages 429–438. IEEE, 2013.
- [15] W. Fan, J. Li, S. Ma, N. Tang, and W. Yu. Towards certain fixes with editing rules and master data. *PVLDB*, 3(1):173–184, 2010.
- [16] J. Feng, H. Xu, S. Mannor, and S. Yan. Robust logistic regression and classification. In *Advances in Neural Information Processing Systems*, pages 253–261, 2014.
- [17] J. Friedman, T. Hastie, and R. Tibshirani. *The elements of statistical learning*, volume 1. Springer series in statistics Springer, Berlin, 2001.
- [18] C. Gokhale, S. Das, A. Doan, J. F. Naughton, N. Rampalli, J. Shavlik, and X. Zhu. Corleone: Hands-off crowdsourcing for entity matching. In *SIGMOD*, 2014.
- [19] A. Guillory, E. Chastain, and J. Bilmes. Active learning as non-convex optimization. In *International Conference on Artificial Intelligence and Statistics*, pages 201–208, 2009.
- [20] J. M. Hellerstein, C. Ré, F. Schoppmann, D. Z. Wang, E. Fratkin, A. Gorajek, K. S. Ng, C. Welton, X. Feng, K. Li, and A. Kumar. The madlib analytics library or MAD skills, the SQL. *PVLDB*, 5(12):1700–1711, 2012.
- [21] M. Jaggi, V. Smith, M. Takác, J. Terhorst, S. Krishnan, T. Hofmann, and M. I. Jordan. Communication-efficient distributed dual coordinate ascent. In *Advances in Neural Information Processing Systems*, pages 3068–3076, 2014.
- [22] S. Kandel, A. Paepcke, J. M. Hellerstein, and J. Heer. Enterprise data analysis and visualization: An interview study. *IEEE Trans. Vis. Comput. Graph.*, 18(12):2917–2926, 2012.
- [23] Z. Khayyat, I. F. Ilyas, A. Jindal, S. Madden, M. Ouzzani, P. Papotti, J.-A. Quiané-Ruiz, N. Tang, and S. Yin. Bigdancing: A system for big data cleansing. 2015.
- [24] H. Köpcke, A. Thor, and E. Rahm. Evaluation of entity resolution approaches on real-world match problems. *PVLDB*, 3(1):484–493, 2010.
- [25] T. Kraska, A. Talwalkar, J. C. Duchi, R. Griffith, M. J. Franklin, and M. I. Jordan. Mlbase: A distributed machine-learning system. In *CIDR 2013, Sixth Biennial Conference on Innovative Data Systems Research, Asilomar, CA, USA, January 6-9, 2013, Online Proceedings*, 2013.
- [26] S. Krishnan, J. Patel, M. J. Franklin, and K. Goldberg. A methodology for learning, analyzing, and mitigating social influence bias in recommender systems. In *Eighth ACM Conference on Recommender Systems, RecSys '14, Foster City, Silicon Valley, CA, USA - October 06 - 10, 2014*, pages 137–144, 2014.
- [27] J. Mahler, S. Krishnan, M. Laskey, S. Sen, A. Murali, B. Kehoe, S. Patil, J. Wang, M. Franklin, P. Abbeel, and K. Y. Goldberg. Learning accurate kinematic control of cable-driven surgical robots using data cleaning and gaussian process regression. In *2014 IEEE International Conference on Automation Science and Engineering, CASE 2014, New Taipei, Taiwan, August 18-22, 2014*, pages 532–539, 2014.
- [28] C. Mayfield, J. Neville, and S. Prabhakar. ERACER: a database approach for statistical inference and data cleaning. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2010, Indianapolis, Indiana, USA, June 6-10, 2010*, pages 75–86, 2010.
- [29] B. Nelson, B. I. Rubinstein, L. Huang, A. D. Joseph, S. J. Lee, S. Rao, and J. Tygar. Query strategies for evading convex-inducing classifiers. *The Journal of Machine Learning Research*, 13(1):1293–1332, 2012.
- [30] A. B. Owen. *Monte Carlo theory, methods and examples*. 2013.
- [31] S. J. Pan and Q. Yang. A survey on transfer learning. *Knowledge and Data Engineering, IEEE Transactions on*, 22(10):1345–1359, 2010.
- [32] H. Park and J. Widom. Crowdfill: collecting structured data from the crowd. In *International Conference on Management of Data, SIGMOD 2014, Snowbird, UT, USA, June 22-27, 2014*, pages 577–588, 2014.
- [33] Z. Qu, P. Richtárik, and T. Zhang. Randomized dual coordinate ascent with arbitrary sampling. *arXiv preprint arXiv:1411.5873*, 2014.
- [34] E. Rahm and H. H. Do. Data cleaning: Problems and current approaches. *IEEE Data Eng. Bull.*, 23(4):3–13, 2000.
- [35] M. J. F. K. G. Sanjay Krishnan, Jiannan Wang and T. Kraska. Stale view cleaning: Getting fresh answers from stale materialized views. *Proceedings of the VLDB Endowment*, 8(12), 2015.
- [36] B. Settles. Active learning literature survey. *University of Wisconsin, Madison*, 52(55-66):11, 2010.
- [37] E. H. Simpson. The interpretation of interaction in contingency tables. *Journal of the Royal Statistical Society. Series B (Methodological)*, pages 238–241, 1951.
- [38] N. Swartz. Gartner warns firms of 'dirty data'. *Information Management Journal*, 41(3), 2007.
- [39] R. Verborgh and M. De Wilde. *Using OpenRefine*. Packt Publishing Ltd, 2013.
- [40] M. Volkovs, F. Chiang, J. Szlichta, and R. J. Miller. Continuous data cleaning. In *Data Engineering (ICDE), 2014 IEEE 30th International Conference on*, pages 244–255. IEEE, 2014.
- [41] M. J. Wainwright, M. I. Jordan, and J. C. Duchi. Privacy aware learning. In *Advances in Neural Information Processing Systems*, pages 1430–1438, 2012.
- [42] J. Wang, T. Kraska, M. J. Franklin, and J. Feng. Crowder: Crowdsourcing entity resolution. *PVLDB*, 5(11):1483–1494, 2012.
- [43] J. Wang, S. Krishnan, M. J. Franklin, K. Goldberg, T. Kraska, and T. Milo. A sample-and-clean framework for fast and accurate query processing on dirty data. In *SIGMOD Conference*, pages 469–480, 2014.
- [44] H. Xiao, T. DE, B. Biggio, D. UNICA, G. Brown, G. Fumera, C. Eckert, I. TUM, and F. Roli. Is feature selection secure against training data poisoning? 2015.
- [45] M. Yakout, A. K. Elmagarmid, J. Neville, M. Ouzzani, and I. F. Ilyas. Guided data repair. *PVLDB*, 2011.
- [46] P. Zhao and T. Zhang. Stochastic optimization with importance sampling. *arXiv preprint arXiv:1401.2753*, 2014.