

ActiveClean: Progressive Data Cleaning For Convex Data Analytics

ABSTRACT

A perennial challenge in data analytics is presence of dirty data in the form of missing, duplicate, incorrect or inconsistent values. The growing popularity of predictive models leads to additional concerns as these models are highly sensitive to systematic corruption. Although errors can be mitigated through data cleaning, it is often very time consuming. Consequently, expensive cleaning is frequently applied progressively; cleaning only as much as necessary to achieve a desired accuracy. However, existing model training methodologies can return misleading results when trained on a mix of dirty and clean data. The key insight of our framework, ActiveClean, is for convex loss models, data cleaning can be directly integrated with model training via Stochastic Gradient Descent allowing for progressive cleaning while preserving provable properties. ActiveClean applies a number of optimizations to improve convergence rates such as importance sampling based on value to the model, avoiding data that is expected to be clean, and batching together updates from already cleaned data. Evaluation on four real-world datasets suggests ActiveClean returns more accurate models than uniform sampling and Active Learning when systematic corruption is sparse. In one set of experiments with 5% simulated error, ActiveClean cleans 55% fewer records to achieve the same accuracy as an Active Learning algorithm.

1. INTRODUCTION

Data are susceptible to various forms of corruption such as missing, incorrect, or inconsistent representations [38]. Dirty data can lead to inaccurate analysis, and a variety of data cleaning techniques have been proposed [35]. The growing popularity of predictive models in data analytics [1,6,12,21] leads to additional challenges in managing dirty data. Predictive models rely on learning relationships between features and labels, and systematic corruption [39] (i.e., corruption that disproportionately affects certain data) can mask or even introduce spurious new relationships. Furthermore, the high dimensionality of these models can amplify small problems [44] resulting in error-prone predictions even when trained on mostly clean data.

To make the systematic corruption problem more concrete, consider a music recommender system in which due to a software bug, all users from Europe have an incorrect age attribute defaulted to “18-24”. A recommendation model trained on this data may spuriously learn a correlation relationship between the “18-24” age group and music liked by European users. A bug, which ostensibly affected only the European users’ record, can affect predictions to all users aged “18-24”. Systematic corruption prior to featurization is not addressed in the robust Machine Learning literature which focuses on the resilience to outliers (i.e., age “150”).

A number of data cleaning frameworks have been recently pro-




Figure 1: ActiveClean is an architecture where data cleaning is integrated with model training in a framework with sampling, model update, and feedback through estimation.

posed to address the problem of corrupted data [5,11,24]. However, data analysts report that data cleaning remains one of the most time consuming steps in the analysis process [3]. Data cleaning can require a significant amount of developer effort in writing software or rules to fix the corruption. Crowdsourcing is an increasingly popular alternative with recent success in missing value filling and entity resolution [5,11,19,33]. However, crowdsourcing comes at the cost of additional latency and the overhead of managing human workers.

When data cleaning is expensive, it is beneficial to apply it *progressively*, where users can inspect early results with only $k \ll N$ records cleaned. Progressive data cleaning allows users to specify a cleaning budget and clean until that budget is reached. Even without a budget, progressive cleaning is useful as it allows users measure the impact of a potentially costly data cleaning operation without cleaning the entire data. However, when applied before predictive modeling, progressive data cleaning poses several methodological problems. Suppose k records are cleaned, but all of the remaining dirty records are retained in the dataset. Training a model on a mixture of dirty and clean data can lead to misleading relationships in even simple scenarios (Figure 2). An alternative is to clean k records and to disregard all of the remaining dirty records (e.g., sampling [43]). While this avoids the mixing problem, accurate model training may require a large amount of training data and k examples may not be enough for a viable model. Finally, both problems are compounded by sparsity, where if corrupted records are uncommon, a random subset of k records may have relatively few examples of corruptions. The errors introduced by these three problems may dominate any gains from data cleaning, leading to unreliable or misleading conclusions about data or model quality.

We propose ActiveClean, a progressive data cleaning framework, that addresses the three methodological challenges: mixing, sampling, and sparsity. The key insight is that an important class of predictive models, called convex loss models (e.g., linear regression and SVMs), are trained by iteratively drawing random samples of data and updating a model [9]. Rather than cleaning before model training, data cleaning can be directly integrated with the sampling and updating training process; preserving provable guarantees such

as convergence and error bounds. In ActiveClean, data are cleaned in small random batches and the model is incrementally updated based on the results. Similar to Active Learning, ActiveClean selects the most valuable records to clean with higher probability, however, it applies a number of optimizations that exploit the data cleaning setting such as avoiding data that is expected to be clean, estimating of the effect of data cleaning for a record, and batching together updates from already cleaned data. This framework is optimized for problems requiring expensive data cleaning.

The ActiveClean architecture (Figure 1) consists of a *detector*, *sampler*, *cleaner*, *updater*, and *estimator*. The cleaner is an existing data cleaning technique (e.g., Entity Resolution), and ActiveClean provides the remaining components to apply this technique progressively. To summarize the contributions in each component:

- **Detector** (Section 4). The detector can apply rules from data quality constraints or adaptively learn which records are dirty to increase the fraction of dirty records sampled.
- **Sampler** (Section 6). We derive an optimal sampling distribution that minimizes the update variance (i.e., how different would the update be if another sample was drawn) which linearly improves an error bound on the convergence rate.
- **Updater** (Section 5). The updater applies a weighted stochastic gradient descent step to the current best model. This update is guaranteed to converge, and for batch size b and iterations T , converges with rate $O(\frac{1}{\sqrt{bT}})$.
- **Estimator** (Section 7) The estimator applies a Taylor Series linearization to decouple changes in different features to use information from error detection to inform estimation.
- The experiments evaluate these components on 4 datasets with real and synthetic corruption (Section 8). The results suggest that indeed ActiveClean is better suited for cleaning and model training when $k \ll N$. For a 5% systematic corruption, ActiveClean cleans 55% fewer records to achieve the same accuracy as an Active Learning algorithm.

2. PROBLEM SETUP

This section describes an example of data cleaning when training predictive models, and formalizes the class of predictive models (convex loss) explored in this work.

2.1 Motivating Scenario

As a motivating example, consider one of the experimental datasets (see Section 8.6.1):

Dollars for Docs [2]. ProPublica collected a dataset of corporate donations to doctors to analyze conflicts of interest. They reported that some doctors received over \$500,000 in travel, meals, and consultation expenses [4]. ProPublica meticulously curated and cleaned a dataset from the Centers for Medicare and Medicaid Services listing nearly 250,000 research donations and aggregated these donations by physician, drug, and pharmaceutical company. We collected the raw unaggregated data and explore whether suspected contributions can be predicted from features. This problem is typical of analysis scenarios based on observational data seen in finance, insurance, and medicine.

The dataset has the following schema:

```
Contribution(pi_specialty, drug_name, device_name,
corporation, amount, dispute, status)
```

`pi_specialty` is a textual attribute describing the specialty of the doctor receiving the donation.

`drug_name` is the branded name of the drug in the research study (null if not a drug).

`device_name` is the branded name of the device in the study (null if not a device).

`corporation` is the name of the pharmaceutical providing the donation.

`amount` is a numerical attribute representing the contribution amount. `dispute` is a Boolean attribute describing whether the research was disputed.

`status` is a string label describing whether the contribution was allowed under the declared research protocol. The goal is to predict disallowed contributions.

However, this dataset is very dirty, and the systematic nature of the corruption can result in an inaccurate model. On the ProPublica website [2], they list numerous types of data problems that had to be cleaned before publishing the data (see Appendix I). For example, the most significant donations were made by large companies whose names were also more often inconsistently represented in the data e.g., “Pfizer Inc.”, “Pfizer Incorporated”, “Pfizer”. In a scenario such as this one, the effect of systematic error can be serious. Duplicate entity representations could artificially reduce the correlation between these entities and suspected contributions. There were nearly 40,000 of the 250,000 records had either entity resolution issues or other inconsistencies in labeling the allowed or disallowed `status`. With the same Support Vector Machine model, the detection rate of suspect donations was 66% in the dirty data and 97% in the clean data (Section 8.6.1).

Progressive Data Cleaning: Cleaning a dataset of 250,000 records can be very time consuming, and it is important for analysts to be able to evaluate the model before all of the data is cleaned. First, suppose k records are cleaned, but all of the remaining dirty records are retained in the dataset. Figure 2 illustrates the dangers of this approach on a very simple dirty dataset and model. The model is a simple linear regression i.e., the best fit line for two variables. One of the variables is systematically corrupted with a translation in the x-axis (Figure 2a). The dirty data is marked in brown and the clean data in orange, and their respective best fit lines are in blue. After cleaning only two of the data points (Figure 2b), the resulting best fit line is in the opposite direction of the true model. This is a well-known phenomenon called Simpson’s paradox, where mixtures of different populations of data can result in spurious relationships [37]. Training models on a mixture of dirty and clean data can lead to unreliable results, where artificial trends introduced by the mixture can be confused for the effects of data cleaning.




Figure 2: (a) Systematic corruption (translation) in one variable where the dirty data is in brown, the clean data is in yellow, and their respective best fit lines are in blue. (b) Applying the same model to mixed data results in an inaccurate model. (c) Likewise, small sample sizes can result in similarly inaccurate models.

An alternative is to avoid the dirty data altogether instead of mixing the two populations. Suppose k records are randomly sampled from the dataset and cleaned. The model is trained only on the cleaned sample of data. This is similar to SampleClean [43], which was proposed to approximate the results of aggregate queries by applying them to a clean sample of data. However, high-dimensional models are highly sensitive to sample size. Figure 2c illustrates

that, even in two dimensions, models trained from small samples can be as incorrect as the mixing solution described before.

In this work, we propose a new methodology to avoid Simpson’s Paradox and the strong dependence on sample size. Instead of mixing dirty and clean data, ActiveClean uses a model trained on the dirty data as an initialization, and then iteratively updates this model using samples of clean data. The intuition is that this algorithm smoothly transitions the model from one population (the dirty data) to another (the clean data), leading to provable guarantees about intermediate results.

2.2 Preliminaries

Convex loss minimization problems (see Friedman, Hastie, and Tibshirani [18] for an introduction) are amenable to incremental stochastic optimization techniques (incremental sample-based updates). This class of problems includes all generalized linear models (including linear and logistic regression), and all variants of support vector machines.

The goal is to learn a vector of model *parameters* θ from training examples $\{(x_i, y_i)\}_{i=1}^N$. θ is learned by minimizing a loss function ϕ (a penalty for getting the prediction wrong) for each training example:

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^N \phi(x_i, y_i, \theta)$$

For example, in a linear regression ϕ is:

$$\phi(x_i, y_i, \theta) = \|\theta^T x_i - y_i\|_2^2$$

Typically, a *regularization* term $r(\theta)$ is added to this problem. $r(\theta)$ penalizes high or low values of feature weights in θ to avoid overfitting to noise in the training examples.

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^N \phi(x_i, y_i, \theta) + r(\theta)$$

In this work, without loss of generality, we will include the regularization as part of the loss function i.e., $\phi(x_i, y_i, \theta)$ includes $r(\theta)$.

2.2.1 Types of Error

There are two standard metrics to quantify inaccuracy in a model:

Model Error. Let θ be the model trained on the dirty data, and let θ^* be the model trained on the same data if it were cleaned. Then the model error is defined as $\|\theta - \theta^*\|$.

Testing Error. Let θ be the model trained on the dirty data, and let θ^* be the model trained on the same data if it were cleaned. Let $T(\theta)$ be the out-of-sample testing error when the dirty model is applied to the clean data, and $T(\theta^*)$ be the test error when the clean model is applied to the clean data. The testing error is defined as $T(\theta^*) - T(\theta)$.

The above *errors* are caused by two underlying problems:

Data Error. is error introduced by training a model on systematically corrupted data.

Sampling Error. Error introduced by sampling i.e., the difference between a model trained on $p\%$ of the data and 100% of the data.

When we use the term *error*, we are referring to **model error**. We will be explicit with other terms when describing data errors (e.g. “data corruption”).

2.3 ActiveClean Problem

The core problem addressed by ActiveClean is incremental model update while progressively cleaning data.

PROBLEM 1 (ACTIVECLEAN PROBLEM). Let R be a dirty relation, $F(r) \mapsto (x, y)$ be a featurization that maps a record $r \in R$ to a feature vector x and label y , ϕ be a convex regular-

ized loss, and $C(r) \mapsto r_{\text{clean}}$ be a cleaning technique that maps a record to its cleaned value. Given these inputs, the ActiveClean problem is to return an estimate $\hat{\theta}$ of the clean model for any limit k on the number of times the data cleaning $C(\cdot)$ can be applied. This estimate should be bounded by a monotonically decreasing function in k (i.e., the expected result of cleaning more data is a more accurate model).

Addressing this problem requires analysis of the loss function ϕ and how ϕ is affected by dirty data. The solution is to integrate data cleaning and model training, where ϕ is simultaneously minimized while the data is cleaned. The tight feedback loop between model training and data cleaning poses several new algorithmic challenges and systems challenges. Algorithmically, we have to re-weight data to avoid biases and the population mixtures described previously. There are also numerous new opportunities for optimizations such as prioritizing data and avoiding data that are expected to be clean.

From a systems perspective, data cleaning and model training happen at very different time scales. When humans are involved, per record latencies for data repair are orders of magnitude larger than the CPU time needed for model training. We can compare recent results in data cleaning to a model training framework like CoCoA implemented on Spark [22]. Per record, BigDancing, a highly optimized automated Spark-based data cleaning system is 15.5x slower than CoCoA¹. Crowd based techniques like Crowd-Fill [33] and CrowdER [42] are over 100,000x slower per record. Consequently, all of the optimizations in ActiveClean are designed to address data cleaning latency (i.e., more progress with fewer cleaned records) rather than optimizing for numerical computation (i.e., process less records).

3. SYSTEM ARCHITECTURE

This section describes the ActiveClean architecture and the basic algorithmic framework. The individual components will be addressed in the subsequent sections.

3.1 Overview

Figure 1 in the introduction overviews the entire framework. The first step of ActiveClean is *initialization*. In this step, there is a dirty relation R , a data cleaning technique $C(\cdot)$, and a dirty model $\theta^{(d)}$ trained on the dirty dataset. Optionally, ActiveClean integrates with dirty data detection rules $D(\cdot)$ which selects the set of likely corrupted records from R . If one is not provided, ActiveClean starts by treating all of the data as dirty and tries to learn a detector as data are cleaned. At initialization, there are two hyperparameters to set, the cleaning budget k and the batch size b (the number of iterations is $T = \frac{k}{b}$). We discuss how to set b and the tradeoffs in setting a larger or smaller b in Section 5.

After initialization, ActiveClean begins the cleaning and model update iterations. The *sampler* selects a sample of dirty data based on the batch size. At this step, ActiveClean can use the detector D to narrow the sample to select only dirty data. Once a sample is selected, the *cleaner* applies $C(\cdot)$ to the dirty sample. Then, after the sample is cleaned, the current model is updated by the *updater*. ActiveClean is initialized with the dirty model, and this model is iteratively updated as more batches are cleaned.

The next two steps in the architecture are feedback steps where the sampling distribution is updated for the next iteration. The *estimator* uses previously cleaned data to estimate the value of data cleaning on new records. This information is used to guide sampling towards more valuable records. After estimation, the detector

¹For CoCoA to reach a precision of 1e-3

$D(\cdot)$ is also updated based on cleaned data. After all of the iterations are complete, the system returns the updated model.

To summarize the architecture in pseudocode:

```

1. Init(dirty_data, cleaned_data, dirty_model, batch,
   iter)
2. For each t in {1, ..., T}
   (a) dirty_sample = Sampler(dirty_data, sample_prob,
   detector, batch)
   (b) clean_sample = Cleaner(dirty_sample)
   (c) current_model = Updater(current_model,
   sample_prob, clean_sample)
   (d) cleaned_data = cleaned_data + clean_sample
   (e) dirty_data = dirty_data - clean_sample
   (f) sample_prob = Estimator(dirty_data, cleaned_data,
   detector)
   (g) detector = DetectorUpdater(detector,
   cleaned_data)
3. Output: current_model

```

Here is an example application of ActiveClean:

EXAMPLE 1. *The analyst first trains an SVM model on the dirty data ignoring the effects of the errors returning a model $\theta^{(d)}$. She decides that she has a budget of cleaning 100 records, and decides to clean the 100 records in batches of 10 (set based on how fast she can clean the data, and how often she wants to see an updated result). She initializes ActiveClean with $\theta^{(d)}$. ActiveClean samples an initial batch of 10 records. She manually cleans those records by merging similar drug names, making corporation names consistent, and fixing incorrect labels. After each batch, the model is updated with the most recent cleaning results $\theta^{(t)}$. The model improves after each iteration. After $t = 10$ of cleaning, the analyst has an accurate model trained with 100 cleaned records but still utilizes the entire dirty data.*

3.2 Challenges and Formalization

We highlight the important components and formalize the research questions explored in this paper.

Detector (Section 4). The first challenge in ActiveClean is dirty data detection. In this step, the detector select a candidate set of dirty records $R_{dirty} \subseteq R$. There are two techniques to do this: (1) an *a priori* case, and (2) an adaptive case. In the *a priori* case, the detector knows which data is dirty in advance. In the adaptive case, the detector learns classifier based on previously cleaned data.

Sampler (Section 6). The sampler draws a sample of records $S_{dirty} \subseteq R_{dirty}$. This is a non-uniform sample where each record r has a sampling probability $p(r)$. We will derive the optimal sampling distribution, and show how the theoretical optimal can be approximated by the next estimator.

Cleaner (User-Specified). Given the sample of the records S_{dirty} , the cleaner applies the user-specified data cleaning $C(\cdot)$. This paper focuses on a record-by-record cleaning model where the function C is applied to a record and produces the clean record:

$$S_{clean} = \{C(r) : \forall r \in S_{dirty}\}$$

This allows for uniform measure of the performance of ActiveClean in terms of model error as a function of sample size. The record-by-record cleaning model is not a fundamental restriction of this approach, and in the extensions (Section A.1), there is a discussion on a compatible “set of records” cleaning model. Consider the case where an analyst finds a dirty record, and is able to fix all records (possibly outside the sample) with same error throughout the dataset efficiently.

Updater (Section 5). The updater updates the model $\theta^{(t)}$ based on the newly cleaned data $F(S_{clean})$ resulting in $\theta^{(t+1)}$. Analyzing the model update procedure as a stochastic gradient descent algorithm will help derive the sampling distribution and estimation.

Estimator (Section 7): The estimator approximates the optimal distribution derived in the Sample step. Based on the change between $F(S_{clean})$ and $F(S_{dirty})$, it directs the next iteration of sampling to select points that will have changes most valuable to the next model update.

4. DETECTION

To maximize the benefit of data cleaning, detection ensures that sampling draws records likely to be dirty.

4.1 Goals

The detector returns two important aspects of a record: (1) whether the record is dirty, and (2) if it is dirty, what is wrong with the record. The sampler can use (1) to select a subset of dirty records to sample at each batch. The estimator can use (2) estimate the value of data cleaning based on other records with the same corruption. ActiveClean supports two types of detectors *a priori* and *adaptive*. In the *a priori* case, there is a way to select the set of dirty records before any cleaning. This case is possible for some types of data errors. In the adaptive case, detection is learned as data is cleaned.

4.2 A Priori Case

For many types of dirtiness such as missing attribute values and constraint violations, it is possible to efficiently enumerate a set of corrupted records and what is wrong with them.

DEFINITION 1 (A PRIORI DETECTION). *Let r be a record in R . An *a priori* detector is a detector that returns a Boolean of whether the record is dirty and a set of columns e_r that are dirty.*

$$D(r) = (\{0, 1\}, e_r)$$

From the set of columns that are dirty, find the corresponding features that are dirty f_r and labels that are dirty l_r .

Here are example use cases of this definition using data cleaning methodologies proposed in the literature.

Constraint-based Repair: One model for detecting errors involves declaring constraints on the database.

Detection. Let Σ be a set of constraints on the relation \mathcal{R} . In the detection step, the detector select a subset of records $\mathcal{R}_{dirty} \subseteq \mathcal{R}$ that violate at least one constraint. The set e_r is the set of columns for each record which have a constraint violation.

EXAMPLE 2. *An example of a constraint on the running example dataset is the status of a contribution can be only “allowed” or “disallowed”. Any other value for status is an error.*

Entity Resolution: Another common data cleaning task is Entity Resolution [19,25,42]. Entity Resolution is the problem of standardizing attributes that represent the same real world entity. A common pattern in Entity Resolution is to split up the operation into two steps: blocking and matching. In blocking, attributes that should be the same are coarsely grouped together. In matching, those coarse groups are resolved to a set of distinct entities.

Detection. Detection for entity resolution problems is the matching step. Let S be a similarity function that takes two records and returns a value in $[0, 1]$ (1 most similar and 0 least similar). For

some threshold t , S defines a similarity relationship between two attributes $r(a)$ and $r'(a)$:

$$r(a) \approx r'(a) : S(r(a), r'(a)) \geq t$$

In the detection step, R_{dirty} is the set of records that have at least one other record in the relation that satisfies $r(a) \approx r(a)'$. The set e_r is the set of attributes of r that have entity resolution problems.

EXAMPLE 3. An example of an Entity Resolution problem is seen in our earlier example about corporation names e.g. “Pfizer Inc.”, “Pfizer Incorporated”, “Pfizer”. Given a similarity relationship $WeightedJaccard(r1, r2) > 0.8$, the detector selects all records that satisfy this condition (their Weighted Jaccard Similarity is greater than 0.8).

4.3 Adaptive Detection

A priori detection is not possible in all cases. The detector also supports adaptive detection where detection is learned from previously cleaned data. Note that this “learning” is distinct from the “learning” at the end of the pipeline. The challenge in formulating this problem is that detector needs to describe how the data is dirty (e.g. e_r in the *a priori* case). The detector achieves this by categorizing the corruption into u classes. These classes are corruption categories that do not necessarily align with features, but every records is classified with at most one category. For example, suppose there are records with outliers and missing values, there are three classes of corruption: outliers, missing values, and both.

When using adaptive detection, the repair step has to clean the data and report to which of the u classes the corrupted record belongs. When an example (x, y) is cleaned, the repair step labels it with one of the clean, 1, 2, ..., $u + 1$ classes (including one for “not dirty”). It is possible that u increases each iteration as more types of dirtiness are discovered. Then, the detection problem reduces to a multiclass classification problem. This problem can be addressed by any multiclass classifier, and we use an all-versus-one SVM in our experiments. Since this classifier is internal to our system, it does not have to be a convex model (i.e., it can be a Decision Tree or Random Forest).

DEFINITION 2 (ADAPTIVE CASE). Select the set of records for which κ gives a positive error classification (i.e., one of the u error classes). After each sample of data is cleaned, the classifier κ is retrained. So the result is:

$$D(r) = (\{1, 0\}, \{1, \dots, u + 1\})$$

Adaptive Detection With Open Refine:

EXAMPLE 4. OpenRefine is a spreadsheet-based tool that allows users to explore and transform data. However, it is limited to cleaning data that can fit in memory on a single computer. Since the cleaning operations are coupled with data exploration, ActiveClean does not know what is dirty in advance (the analyst may discover new errors as she cleans).

Suppose the analyst wants to use OpenRefine to clean the running example dataset with ActiveClean. She takes a sample of data from the entire dataset and uses the tool to discover errors. For example, she finds that some drugs are incorrectly classified as both drugs and devices. She then clears the device attribute for all records that have the drug name in question. Every time she makes a batch data transformation (i.e., cleaning the device attribute), ActiveClean can list the set of records that have changed. Each transformation becomes an error class, and the records that have changed records become positive training examples for a classifier to guide future samples.

5. UPDATE

Before discussing sampling, this section discusses how to update a model. The sampling framework will naturally follow from this update procedures. This section assumes that the detector in the previous section has selected a set of candidate records R_{dirty} and the sampler has sampled from this set of candidate records. This section show that this model update procedure can be interpreted as a Stochastic Gradient Descent (SGD) algorithm, which gives a theoretical framework to analyze convergence and bound the error at each step.

5.1 Update Problem

Let $S_{dirty} \subseteq R_{dirty}$ be a random sample of data from the results of the detection in the previous section, and let $p(r)$ be the sampling probability of each record. The cleaner applies the data cleaning technique to the sample resulting in S_{clean} and clean features and labels $(X^{(c)}, Y^{(c)})$. In the model update problem, the updater has to update the dirty model $\theta^{(d)}$ based on the clean sample and return θ^{new} . The goal is that these updates should minimize the error of the updated model and the true model $\theta^{(c)}$ (if entire data is cleaned):

$$error(\theta^{new}) = \|\theta^{new} - \theta^{(c)}\|$$

5.2 Geometric Interpretation

The update algorithm intuitively follows from the convex geometry of the problem. Consider this problem in one dimension (i.e., the parameter θ is a scalar value), so then the goal is to find the minimum point (θ) of a curve $l(\theta)$. The consequence of dirty data is that the wrong loss function is optimized. Figure 3A illustrates the consequence of this optimization. The brown dotted line shows the loss function on the dirty data. Optimizing this loss function finds θ that at the minimum point. However, the true loss function (w.r.t to the clean data) is in blue. This optimal value on the dirty data is a suboptimal point on clean curve.




Figure 3: (A) A model trained on dirty data can be thought of as a sub-optimal point w.r.t to the clean data. (B) The gradient gives us the direction to move the suboptimal model to approach the true optimum.

The optimal clean model $\theta^{(c)}$ is visualized as a yellow star. The first question is which direction to update θ (i.e., left or right). For this class of models, given a suboptimal point, the direction to the global optimum is the gradient of the loss function. The gradient is a d -dimensional vector function of the current model θ and the clean data. Given this direction, the updater needs to update $\theta^{(d)}$ some distance γ (Figure 3b):

$$\theta^{new} \leftarrow \theta^{(d)} - \gamma \cdot \nabla \phi(\theta^{(d)})$$

At the optimal point, the magnitude of the gradient will be zero. So intuitively, this approach iteratively moves the model downhill, or corrects the dirty model until the desired accuracy is reached.

However, the gradient depends on all of the clean data which is not available and the updater will have to approximate this gradient from a sample. The intuition, formalized in Section 5.5, is that if

the gradients are on average in the right direction, the algorithm is guaranteed to converge with bounds on the convergence rate.

5.3 Average Gradient From a Sample

To derive a sample-based update rule, the most important property is that sums commute with derivatives and gradients. The convex loss class of models are sums of losses, so given the current best model θ , the gradient $g^*(\theta)$ is:

$$g^*(\theta) = \nabla \phi(\theta) = \frac{1}{N} \sum_i^N \nabla \phi(x_i^{(c)}, y_i^{(c)}, \theta)$$

Therefore, the gradient can be estimated from a sample by taking the gradient w.r.t each record, and re-weighting the average by their respective sampling probabilities. Let S be a sample of data, where each $i \in S$ is drawn with probability $p(i)$:

$$g^*(\theta) \approx g_S(\theta) = \frac{1}{n |S|} \sum_{i \in S} \frac{1}{p(i)} \nabla \phi(x_i^{(c)}, y_i^{(c)}, \theta)$$

Now adding iteration, for every batch of data cleaned t , the update to the current best model estimate is:

$$\theta^{(t+1)} \leftarrow \theta^{(t)} - \gamma \cdot g_S(\theta^{(t)})$$

The detection in the previous step adds a small complication, since the sample S_{dirty} is not representative of all of the data. This requires a compensation for this bias by averaging this estimate with the gradient of the complement:

$$g_C(\theta) = \frac{1}{|R - R_{dirty}|} \sum_{i \in R - R_{dirty}} \nabla \phi(x_i^{(c)}, y_i^{(c)}, \theta)$$

Then, for weights α, β (discussed in Section 5.5):

$$g(\theta) = \alpha \cdot g_C(\theta) + \beta \cdot g_S(\theta)$$

Finally, adding in the iteration, and at each iteration t , the update becomes:

$$\theta^{(t+1)} \leftarrow \theta^{(t)} - \gamma \cdot g(\theta^{(t)}) \blacksquare$$

5.4 Model Update Algorithm

To summarize, the algorithm is initialized with $\theta^{(1)} = \theta^{(d)}$ which is the dirty model. At each iteration $t = \{1, \dots, T\}$, the cleaning is applied to a batch of data b selected from the set of candidate dirty records R_{dirty} . Then, an average gradient is estimated from the cleaned batch and the model is updated. Iteration continues until $k = T \cdot b$ records are cleaned.

1. Calculate the gradient over the sample of clean data and call the result $g_S(\theta^{(t)})$
2. Calculate the average gradient over all the data in $R_{clean} = R - R_{dirty}$, and call the result $g_C(\theta^{(t)})$
3. Apply the following update rule:

$$\theta^{(t+1)} \leftarrow \theta^{(t)} - \lambda \cdot (\alpha \cdot g_S(\theta^{(t)}) + \beta \cdot g_C(\theta^{(t)}))$$

To summarize the parameters $b, \lambda, \alpha, \beta$.

Batch Size b : The batch size b controls the frequency of iteration. Larger batches provide a more accurate estimate of the gradient at each iteration but has less frequent model updates.

Step Size γ : γ controls how far should to travel in the gradient direction.

Weights α, β : α, β are the proportions to combine $g_S(\theta)$ and $g_C(\theta)$.

5.5 Analysis with Stochastic Gradient Descent

This update policy can be formalized as a class of very well studied algorithms called Stochastic Gradient Descent. This provides a theoretical framework to understand and analyze the update rule, bound the error, and choose points to clean. Mini-batch stochastic gradient descent (SGD) is an algorithm for finding the optimal

value given the convex loss and data. In mini-batch SGD, random subsets of data are selected at each iteration and the average gradient is computed for every batch. The key difference is that in traditional SGD there is no notion of dirty and clean data.

ActiveClean as Lazy SGD: This update method is a variant of SGD that lazily materializes the clean value. As data is sampled at each iteration, data is cleaned when needed by the optimization. It is well known that even for arbitrary initialization SGD makes significant progress in less than one epoch (a pass through the entire dataset) [10]. Furthermore in this setting, the dirty model can be much more accurate than an arbitrary initialization; leading to highly accurate models without cleaning the entire data.

Deriving α and β : The first problem is to choose α and β such that the estimate of the gradient is unbiased. SGD will converge when the estimate is unbiased and the step-size is chosen appropriately. The batch S_{dirty} is drawn only from R_{dirty} . Since the sizes of R_{dirty} and its complement are known, it follows that the gradient over the already clean data g_C and the recently cleaned data g_S can be combined as follows:

$$g(\theta^t) = \frac{|R_{dirty}| \cdot g_S + |R_{clean}| \cdot g_C}{|R|}$$

Therefore,

$$\alpha = \frac{|R_{clean}|}{|R|}, \beta = \frac{|R_{dirty}|}{|R|d}$$

LEMMA 1. *The gradient estimate $g(\theta)$ is unbiased if g_S is an unbiased estimate of:*

$$\frac{1}{|R_{dirty}|} \sum g_i(\theta)$$

PROOF SKETCH. This result follows directly from the linearity of expectation (See Appendix B). \square

Setting γ : There is extensive literature in machine learning for choosing the step size γ appropriately. γ can be set either to be a constant or decayed over time. Many machine learning frameworks (e.g., MLlib, Sci-kit Learn, Vowpal Wabbit) automatically set learning rates or provide different learning scheduling frameworks. In the experiments, we use a technique called inverse scaling where there is a parameter $\gamma_0 = 0.1$, and at each iteration it decays to $\gamma_t = \frac{\gamma_0}{|S|t}$.

Convergence: Convergence properties of batch SGD formulations have been well studied [13]. From the preceding analysis, the gradient estimate is unbiased and the step size is appropriately chosen. Following from these two points, convergence is ensured:

PROPOSITION 1. *For an appropriately chosen learning rate γ_t , batch stochastic gradient descent will converge if $\mathbb{E}(g_S) = g^*$.*

Convergence Rate: The convergence rates of SGD are also well analyzed [9,13,45]. This gives a bound on the error of intermediate models and the expected number of steps before achieving a model within a certain error.

PROPOSITION 2. *For a general convex loss, a batch size b , and T iterations, the convergence rate is bounded by $O(\frac{\sigma^2}{\sqrt{bT}})$. σ^2 is the variance in the estimate of the gradient at each iteration:*

$$\mathbb{E}(\|g_S - g^*\|^2)$$

Setting the batch size: The batch size should be set by the user to have the desired properties. Larger batches will take longer to clean and will make more progress towards the clean model, but will have less frequent model updates. On the other hand, smaller batches are cleaned faster and have more frequent model updates

but make less progress. The overheads introduced by ActiveClean are more evident at smaller batch sizes. There are diminishing returns to increasing the batch size $O(\frac{1}{\sqrt{b}})$. Empirically, in the experiments, batch sizes of 50 converge the fastest. If a data cleaning technique requires a larger batch size than this, i.e., data cleaning is fast enough that the iteration overhead is significant compared to cleaning 50 records, ActiveClean can apply the updates in smaller batches. For example, the batch size set by the user might be $b = 1000$, but the model updates after every 50 records are cleaned. This dissociates the batching requirements of SGD and the batching requirements of the data cleaning technique.

Non-convex losses: (See Appendix A.2 for a note on such losses.)

6. SAMPLING

In the previous section, the model update received a sample with probabilities $p(r)$. This section provides a derivation for an optimal sampling problem that directly follows from the analysis of the update rule via SGD. It will turn out that the solution to the optimal sampling problem is not realizable in practice (as it depends on knowing the cleaned value), and this problem will be addressed with an approximation in the next section.

6.1 Goals and Challenges

In the Machine Learning and Optimization literature, SGD algorithms are optimized to avoid scanning the entire data. Uniform sampling is cheap so it is the preferred solution. However, data cleaning costs can be many orders of magnitude higher than model training. As a result, uniform sampling may not be the most efficient option. ActiveClean can sacrifice computational overhead by precomputing some results over the entire data for savings during the data cleaning phase. This problem is formulated as an optimal sampling problem to compute the sampling probabilities $p(r)$ that maximize the convergence rate.

6.2 Optimal Sampling Problem

Recall that the convergence rate of an SGD algorithm is bounded by σ^2 which is the variance of the gradient. Intuitively, the variance measures how accurately the gradient is estimated from a uniform sample. Other sampling distributions, while preserving the sample expected value, may have a lower variance. Thus, the optimal sampling problem is defined as a search over sampling distributions to find the minimum variance sampling distribution.

DEFINITION 3 (OPTIMAL SAMPLING PROBLEM). *Given a set of candidate dirty data R_{dirty} , $\forall r \in R_{dirty}$ find sampling probabilities $p(r)$ such that over all samples S of size k it minimizes:*

$$\mathbb{E}(\|g_S - g^*\|^2)$$

To construct these sampling probabilities, first consider the following lemma about importance sampling. This lemma describes the optimal distribution over a set of scalars:

LEMMA 2. *Given a set of real numbers $A = \{a_1, \dots, a_n\}$, let \hat{A} be a sample with replacement of A of size k . If μ is the mean \hat{A} , the sampling distribution that minimizes the variance of μ , i.e., the expected square error, is $p(a_i) \propto a_i$.*

PROOF SKETCH. This proof follows from [31], as it is a straightforward importance sampling result. We include the proof in the appendix (Section D) \square

Lemma 2 shows that when estimating a mean of numbers with sampling, the distribution with optimal variance is sampling proportionally to the values. This insight leads to a direct higher-

dimensional generalization, where at iteration t the optimal distribution over records in R_{dirty} is probabilities proportional to:

$$p_i \propto \|\nabla\phi(x_i^{(c)}, y_i^{(c)}, \theta^{(t)})\|$$

However, in this case, it leads to a chicken-and-egg problem. The optimal sampling distribution requires knowing $(x_i^{(c)}, y_i^{(c)})$, however, cleaning is required to know those values. In the next section, the estimator will approximate this distribution by estimating the cleaned value with previously cleaned data. Since the model update can work with *any* distribution, convergence is guaranteed no matter how inaccurate this approximation is. However, a better approximation will lead to an improved convergence rate.

7. ESTIMATION

This section makes the sampling result of the previous section practical by approximating the cleaned values.

7.1 Challenges and Goal

The optimal sampling distribution is dependent on a value that is not known without data cleaning $\nabla\phi(x_i^{(c)}, y_i^{(c)}, \theta^{(t)})$. One way to approximate this distribution is to learn a function $e(\cdot)$ via regression based on previously cleaned data. This is a high-dimensional regression problem which may have to learn a very complicated relationship between dirty and clean data. The biggest challenge with such an estimator is the cold start problem, where if given a small amount of cleaned data, the estimator will be inaccurate. ActiveClean should be optimized to make as much progress as possible in the early iterations so this technique may not work. The estimator takes an alternative approach where it exploits information from the detector to produce an estimate for groups of similarly corrupted records.

7.2 Estimation For A Priori Detection

EXAMPLE 5. *Suppose records from running example dataset are corrupted with both entity resolution problems, missing data, and constraint violations. Each training example will have a set of corrupted features (e.g., $\{1, 2, 6\}$, $\{1, 2, 15\}$).*

Suppose that the cleaner has just cleaned the records r_1 and r_2 represented as tuples with their corrupted feature set: $(r_1, \{1, 2, 3\})$, $(r_2, \{1, 2, 6\})$. Then, given a new record $(r_3, \{1, 2, 3, 6\})$. The estimator should be able to use the cleaning results from r_1, r_2 to estimate the gradient in r_3 .

If most of the features are correct, it would seem like the gradient is only incorrect in one or two of its components. The problem is that the gradient $\nabla\phi(\cdot)$ can be a very non-linear function of the features that couple features together. For example, the gradient for linear regression is:

$$\nabla\phi(x, y, \theta) = (\theta^T x - y)x$$

It is not possible to isolate the effect of a change of one feature on the gradient. Even if one of the features is corrupted, all of the gradient components will be incorrect.

7.2.1 Error Decoupling

To address this problem, the gradient can be approximated in a way that the effects of dirty features on the gradient are decoupled. Recall, in the *a priori* detection problem, that associated with each $r \in R_{dirty}$ is a set of errors f_r, l_r which is a set that identifies a set of corrupted features and labels. This property can be used to construct a coarse estimate of the clean value. The main idea is to calculate average changes for each feature, then given an uncleaned (but dirty) record, add these average changes to correct the gradient.

To formalize this intuition, instead of computing the actual gradient with respect to the true clean values, compute the conditional

expectation given that a set of features and labels f_r, l_r are corrupted:

$$p_i \propto \mathbb{E}(\nabla\phi(x_i^{(c)}, y_i^{(c)}, \theta^{(t)}) \mid f_r, l_r)$$

Corrupted features defined as that:

$$i \notin f_r \implies x^{(c)}[i] - x^{(d)}[i] = 0$$

$$i \notin l_r \implies y^{(c)}[i] - y^{(d)}[i] = 0$$

The needed approximation represents a linearization of the errors, and the resulting approximation will be of the form:

$$p_r \propto \|\nabla\phi(x, y, \theta^{(t)}) + M_x \cdot \Delta_{rx} + M_y \cdot \Delta_{ry}\|$$

where M_x, M_y are matrices and Δ_{rx} and Δ_{ry} are average change vectors for the corrupted features in r . Without this approximation, calculating the expected value conditioned on f_r, l_r would require conditioning on all the combinatorial possibilities.

7.2.2 Deriving M_x, M_y

If d is the dirty value and c is the clean value, the Taylor series approximation for a function f is given as follows:

$$f(c) = f(d) + f'(d) \cdot (d - c) + \dots$$

Ignoring the higher order terms, the linear term $f'(d) \cdot (d - c)$ is a linear function in each feature and label. Then, taking expected values, it follows that:

$$\approx \nabla\phi(x, y, \theta) + M_x \cdot \mathbb{E}(\Delta x) + M_y \cdot \mathbb{E}(\Delta y)$$

where $M_x = \frac{\partial}{\partial X} \nabla\phi$ and $M_y = \frac{\partial}{\partial Y} \nabla\phi$ (See Appendix E for derivation). Recall that the feature space is d dimensional and label space is l dimensional. Then, M_x is an $d \times d$ matrix, and M_y is a $d \times l$ matrix. Both of these matrices are computed for each record (see Appendix F for an example derivation). Δx is a d dimensional vector where each component represents a change in that feature and Δy is an l dimensional vector that represents the change in each of the labels.

7.2.3 More Accurate Early Error Estimates

Linearization over avoids amplifying estimation error. Consider the linear regression gradient:

$$\nabla\phi(x, y, \theta) = (\theta^T x - y)x$$

This can be rewritten as a vector in each component:

$$g[i] = \sum_i x[i]^2 - x[i]y + \sum_{j \neq i} \theta[j]x[j]$$

This function is already mostly linear in x except for the one quadratic term. However, this one quadratic term has potential to amplify errors. Consider two expressions:

$$f(x + \epsilon) = (x + \epsilon)^2 = x^2 + 2x\epsilon + \epsilon^2$$

$$f(x + \epsilon) \approx f(x) + f'(x)(\epsilon) = x^2 + 2x\epsilon$$

The only difference between the two estimates is the quadratic ϵ^2 , if ϵ is highly uncertain random variable then the quadratic dominates. If this variance is large, the Taylor estimate avoids amplifying this error. Of course, this is at the tradeoff of some additional bias since the true function is non-linear. We evaluate this linearization in Section 8.5 against alternatives, and find that indeed it provides more accurate estimates for a small number of samples cleaned. When the number of cleaned samples is large the alternative techniques are comparable or even slightly better.

7.2.4 Maintaining Decoupled Averages

This linearization allows ActiveClean to maintain per feature (or label) average changes and use these changes to center the optimal sampling distribution around the expected clean value. To estimate $\mathbb{E}(\Delta x)$ and $\mathbb{E}(\Delta y)$, consider the following lemma:

LEMMA 3 (SINGLE FEATURE). *For a feature i , we average all $j = \{1, \dots, K\}$ records cleaned that have an error for that fea-*

ture, weighted by their sampling probability:

$$\bar{\Delta}_{xi} = \frac{1}{NK} \sum_{j=1}^K (x^{(d)}[i] - x^{(c)}[i]) \times \frac{1}{p(j)}$$

Similarly, for a label i :

$$\bar{\Delta}_{yi} = \frac{1}{NK} \sum_{j=1}^K (y^{(d)}[i] - y^{(c)}[i]) \times \frac{1}{p(j)}$$

Each $\bar{\Delta}_{xi}$ and $\bar{\Delta}_{yi}$ represents an average change in a single feature. A single vector can represent the necessary changes to apply to a record r :

LEMMA 4 (DELTA VECTOR). *For a record r , the set of corrupted features is f_r, l_r . Then, each record r has a d -dimensional vector Δ_{rx} which is constructed as follows:*

$$\Delta_{rx}[i] = \begin{cases} 0 & i \notin f_r \\ \bar{\Delta}_{xi} & i \in f_r \end{cases}$$

Each record r also has an l -dimensional vector Δ_{ry} which is constructed as follows:

$$\Delta_{ry}[i] = \begin{cases} 0 & i \notin l_r \\ \bar{\Delta}_{yi} & i \in l_r \end{cases}$$

Finally, the result is:

$$p_r \propto \|\nabla\phi(x, y, \theta^{(t)}) + M_x \cdot \Delta_{rx} + M_y \cdot \Delta_{ry}\| \blacksquare$$

7.3 Estimation For Adaptive Case

A similar procedure holds in the adaptive setting, however, it requires reformulation of ‘‘similarly corrupted’’. Here, ActiveClean uses u corruption classes provided by the detector. Instead of conditioning on the features that are corrupted, the estimator conditions on the classes. So for each error class, it computes a Δ_{ux} and Δ_{uy} . These are the average change in the features given that class and the average change in labels given that class.

$$p_{r,u} \propto \|\nabla\phi(x, y, \theta^{(t)}) + M_x \cdot \Delta_{ux} + M_y \cdot \Delta_{uy}\| \blacksquare$$

8. EXPERIMENTS

First, the experiments evaluate how various types corrupted data benefit from data cleaning. Next, the experiments explore different prioritization and model update schemes for progressive data cleaning. Finally, ActiveClean is evaluated end-to-end in a number of real-world data cleaning scenarios.

8.1 Experimental Setup and Notation

The main metric for evaluation is a relative measure of the trained model and the model if all of the data is cleaned.

Relative Model Error. Let θ be the model trained on the dirty data, and let θ^* be the model trained on the same data if it was cleaned. Then the model error is defined as $\frac{\|\theta - \theta^*\|}{\|\theta^*\|}$.

8.1.1 Scenarios

Income Classification (Adult): In this dataset of 45,552 records, the task is to predict the income bracket (binary) from 12 numerical and categorical covariates with an SVM classifier.

Seizure Classification (EEG): In this dataset, the task is to predict the onset of a seizure (binary) from 15 numerical covariates with a thresholded Linear Regression. There are 14980 data points in this dataset. This classification task is inherently hard with an accuracy on completely clean data of only 65%.

Handwriting Recognition (MNIST)²: In this dataset, the task is

²http://ufldl.stanford.edu/wiki/index.php/Using_the_MNIST_Dataset




Figure 4: (a) Robust techniques and discarding data work when corrupted data are random and look atypical. (b) Data cleaning can provide reliable performance in both the systematically corrupted setting and randomly corrupted setting.

to classify 60,000 images of handwritten images into 10 categories with an one-to-all multiclass SVM classifier. The unique part of this dataset is the featurized data consists of a 784 dimensional vector which includes edge detectors and raw image patches.

Dollars For Docs: The dataset has 240,089 records with 5 textual attributes and one numerical attribute. The dataset is featurized with bag-of-words featurization model for the textual attributes which resulted in a 2021 dimensional feature vector, and a binary SVM is used to classify the status of the medical donations.

8.1.2 Compared Algorithms

Here are the alternative methodologies evaluated in the experiments:

Robust Logistic Regression [17]. Feng et al. proposed a variant of logistic regression that is robust to outliers. We chose this algorithm because it is a robust extension of the convex regularized loss model, leading to a better apples-to-apples comparison between the techniques. (See details in Appendix H.1)

Discarding Dirty Data. As a baseline, dirty data is discarded.

SampleClean (SC) [43]. SampleClean takes a sample of data, applies data cleaning, and then trains a model to completion on the sample.

Active Learning (AL) [20]. An Active Learning algorithm that integrates with stochastic optimization (See details in Appendix H.2).

ActiveClean Oracle (AC+O): In ActiveClean Oracle, instead of an estimation step, the true clean value is used to evaluate the theoretical ideal performance of ActiveClean.

8.2 Does Data Cleaning Matter?

The first experiment evaluates the benefits of data cleaning on 2 of the example datasets (EEG and Adult). This is done without sampling to understand which types of data corruption are amenable to data cleaning and which are better suited for robust statistical techniques. The experiment compares 4 schemes: (1) full data cleaning, (2) baseline of no cleaning, (3) discarding the dirty data, and (4) robust logistic regression. We corrupted 5% of the training examples in each dataset in two different ways:

Random Corruption: Simulated high-magnitude random outliers. 5% of the examples are selected at random and a random feature is replaced with 3 times the highest feature value.

Systematic Corruption: Simulated innocuous looking (but still incorrect) systematic corruption. The model is trained on the clean data, and the three most important features (highest weighted) are identified. The examples by each these features and the top examples $\frac{2}{3}$ % of the examples are corrupted with the mean value for that feature. It is important to note the some examples can have multiple corrupted features.

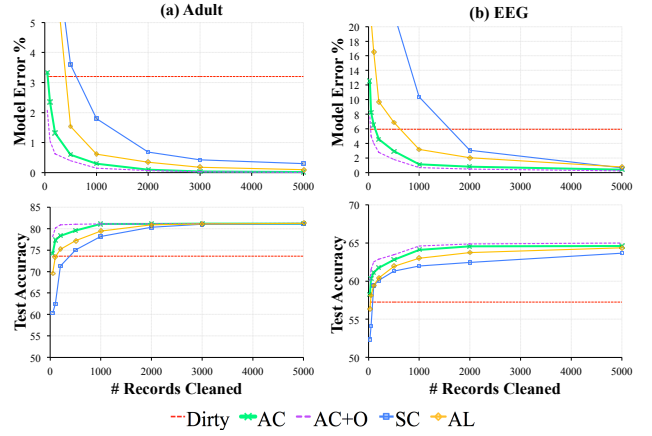


Figure 5: The relative model error as a function of the number of examples cleaned. ActiveClean converges with a smaller sample size to the true result in comparison to Active Learning and SampleClean.

Figure 4 shows the test accuracy for models trained on both types of data with the different techniques. The robust method performs well on the random high-magnitude outliers with only a 2.0% reduction in clean test accuracy for EEG and 2.5% reduction for Adult. In the random setting, discarding dirty data also performs relatively well. However, the robust method falters on the systematic corruption with a 9.1% reduction in clean test accuracy for EEG and 10.5% reduction for Adult. The problem is that without cleaning, there is now way to know if the corruption is random or systematic and when to trust a robust method. While data cleaning requires more effort, it provides benefits in both settings. In the remaining experiments, unless otherwise noted, the experiments use systematic corruption.

Summary: A 5% systematic corruption can introduce a 10% reduction in test accuracy even when using a robust method.

8.3 ActiveClean: A Priori Detection

The next set of experiments evaluate different approaches to cleaning a sample of data compared to ActiveClean using *a priori* detection. *A priori* detection assumes that all of the corrupted records are known in advance but their clean values are unknown.

8.3.1 Active Learning and SampleClean

The next experiment evaluates the samples-to-error tradeoff between three alternative algorithms: ActiveClean (AC), SampleClean, Active Learning, and ActiveClean + Oracle (AC+O). Figure 5 shows the model error and test accuracy as a function of the number of cleaned records. In terms of model error, ActiveClean gives its largest benefits for small sample sizes. For 500 cleaned records of the Adult dataset, ActiveClean has 6.1x less error than SampleClean and 2.1x less error than Active Learning. For 500 cleaned records of the EEG dataset, ActiveClean has 9.6x less error than SampleClean and 2.4x less error than Active Learning. Both Active Learning and ActiveClean benefit from the initialization with the dirty model as they do not retrain their models from scratch, and ActiveClean improves on this performance with detection and error estimation. These gains also correlate well to improvements in test accuracy. For example, to achieve a test accuracy of 80% on the Adult dataset, ActiveClean cleans 500 fewer records than Active Learning.

Summary: ActiveClean with a priori detection returns results that are more than 6x more accurate than SampleClean and 2x more accurate than Active Learning for cleaning 500 records.

8.3.2 Source of Improvements

Figure 5 compares the performance of ActiveClean with and without various optimizations at 500 records cleaned point. This is a vertical slice of the plots in the previous experiments. ActiveClean without detection is denoted as (AC-D) (that is at each iteration we sample from the entire dirty data), and ActiveClean without detection and importance sampling is denoted as (AC-D-I). Figure 6 plots the relative error of the alternatives and ActiveClean with and without the optimizations. Without detection (AC-D), ActiveClean is still more accurate than Active Learning. Removing the importance sampling, ActiveClean is slightly worse than Active Learning on the Adult dataset but is comparable on the EEG dataset. Active Learning is not always guaranteed to outperform a passive approach and on a harder dataset, namely the EEG dataset, ActiveClean with uniform sampling performs just as well.

Summary: Both a priori detection and non-uniform sampling significantly contribute to the gains over Active Learning.




Figure 6: -D denotes no detection, and -D-I denotes no detection and no importance sampling. Both optimizations significantly help ActiveClean outperform SampleClean and Active Learning.

8.3.3 Mixing Dirty and Clean Data

Training a model on mixed data is an unreliable methodology lacking the same guarantees as Active Learning or SampleClean even in the simplest of cases. For thoroughness, the next experiments include the model error as a function of records cleaned in comparison to ActiveClean. Figure 7 plots the same curves as the previous experiment comparing ActiveClean, Active Learning, and two mixed data algorithms. PC randomly samples data, clean, and writes-back the cleaned data. PC+D randomly samples data from using the dirty data detector, cleans, and writes-back the cleaned data. For these errors PC and PC+D give reasonable results (not always guaranteed), but ActiveClean converges faster. This is because ActiveClean tunes the weighting when averaging dirty and clean data into the gradient.




Figure 7: The relative model error as a function of the number of examples cleaned. ActiveClean converges with a smaller sample size to the true result in comparison to partial cleaning (PC, PC+D).

Summary: ActiveClean converges faster than mixing dirty and clean data since it reweights data based on the fraction that is dirty and

clean. Partial cleaning is not guaranteed to give sensible results.

8.3.4 Corruption Rate

The next experiment explores how much of the performance is due to the initialization with the dirty model (i.e., SampleClean trains a model from “scratch”). Figure 8 varies the systematic corruption rate and plots the number of records cleaned to achieve 1% relative error for SampleClean and ActiveClean. SampleClean does not use the dirty data and thus its error is essentially governed by the Central Limit Theorem. SampleClean outperforms ActiveClean only when corruptions are very severe (45% in Adult and nearly 60% in EEG). When the initialization with the dirty model is inaccurate, ActiveClean does not perform as well.

Summary: SampleClean is beneficial in comparison to ActiveClean when corruption rates exceed 45%.




Figure 8: ActiveClean performs well until the corruption is so severe that the dirty model is not a good initialization.

8.4 ActiveClean: Adaptive Detection

This experiment explores how the results of the previous experiment change when using an adaptive detector instead of the *a priori* detector. Recall, in the systematic corruption, 3 of the most informative features were corrupted, thus we group these problems into 6 classes. We use an all-versus-one SVM to learn the categorization.

8.4.1 Basic Performance

Figure 9 overlays the convergence plots in the previous experiments with a curve (denoted by AC+C) that represents ActiveClean using a classifier instead of the *a priori* detection. Initially ActiveClean is comparable to Active Learning; however, as the classifier becomes more effective the detection improves the performance. Over both datasets, at the 500 records point on the curve, adaptive ActiveClean has a 30% higher model error compared to *a priori* ActiveClean. At 1000 records point on the curve, adaptive ActiveClean has about 10% higher error.

Summary: For 500 records cleaned, adaptive ActiveClean has a 30% higher model error compared to a priori ActiveClean, but still outperforms Active Learning and SampleClean.

8.4.2 Classifiable Errors

The adaptive case depends on being able to predict corrupted records. For example, random corruption that look like other data may be hard to learn. As corruption becomes more random, the classifier becomes increasingly erroneous. The next experiment explores making the systematic corruption more random. Instead of selecting the highest valued records for the most valuable features, we corrupt random records with probability p . We compare these results to AC-D where we do not have a detector at all at one vertical slice of the previous plot (cleaning 1000 records). Figure 10a plots the relative error reduction using a classifier. When the corruption is about 50% random then there is a break even point where no detection is better. This is because the classifier is imperfect and misclassifies some data points incorrectly as cleaned.




Figure 9: Even with a classifier ActiveClean converges faster than Active Learning and SampleClean.

Summary: When errors are increasingly random (50% random) and cannot be accurately classified, adaptive detection provides no benefit over no detection.




Figure 10: (a) Data corruptions that are less random are easier to classify, and lead to more significant reductions in relative model error. (b) The Taylor series approximation gives more accurate estimates when the amount of cleaned data is small.

8.5 Estimation

The next experiment compares estimation techniques: (1) “linear regression” train a linear regression model that predicts the clean gradient as a function of the dirty gradient, (2) “average gradient” which does not use the detection to inform how to apply the estimate, (3) “average feature change” uses detection but no linearization, and (4) the Taylor series linear approximation. Figure 10b measures how accurately each estimation technique estimates the gradient as a function of the number of cleaned records on the EEG dataset.

Estimation error is measured using the relative L2 error with the true gradient. The Taylor series approximation proposed gives more accurate for small cleaning sizes, confirming the analysis in Section 7.2.3. Linear regression and the average feature change technique do eventually perform comparably but only after cleaning much more data.

Summary: Linearized gradient estimates are more accurate when estimated from small samples.

8.6 Real World Scenarios

The next set of experiments evaluate ActiveClean in two real world scenarios, one demonstrating the *a priori* case and one for the adaptive detection case.

8.6.1 A Priori: Constraint Cleaning

The first scenario explores the Dollars for Docs dataset published by ProPublica described throughout the paper. To run this experiment, the entire dataset was cleaned up front, and simulated sampling from the dirty data and cleaning by looking up the value in the cleaned data (see Appendix I for constraints, errors, and cleaning methodology). Figure 11a shows that ActiveClean converges faster than Active Learning and SampleClean. To achieve a 4% relative error (i.e., a 75% error reduction from the dirty model), ActiveClean cleans 40000 fewer records than Active Learning. Also, for

10000 records cleaned, ActiveClean has nearly an order of magnitude smaller error than SampleClean.

Figure 11b shows the detection rate (fraction of disallowed research contributions identified) of the classifier as a function of the number of records cleaned. On the dirty data, we can only correctly classify 66% of the suspected examples (88% overall accuracy due to a class imbalance). On the cleaned data, this classifier is nearly perfect with a 97% true positive rate (98% overall accuracy). ActiveClean converges to the cleaned accuracy faster than the alternatives with a classifier of 92% true positive rate for only 10000 records cleaned.

Summary: To achieve an 80% detection rate, ActiveClean cleans nearly 10x less records than Active Learning.




Figure 11: (A) The relative model error as a function of the number of cleaned records. (B) The true positive rate as a function of the number of cleaned records.

8.6.2 Adaptive: Replacing Corrupted Data

The next experiment explores the MNIST handwritten digit recognition dataset with a MATLAB image processing pipeline. In this scenario, the analyst must inspect a potentially corrupted image and replace it with a higher quality one. The MNIST dataset consists of 64x64 grayscale images. There are two types of simulated corruptions: (1) 5x5 block removal where a random 5x5 block is removed from the image by setting its pixel values to 0, and (2) Fuzzy where a 4x4 moving average patch is applied over the entire image. These corruptions are applied to a random 5% of the images, and mimic the random (Fuzzy) vs. systematic corruption (5x5 removal) studied in the previous experiments. The adaptive detector uses a 10 class classifier (one for each digit) to detect the corruption.

Figure 12 shows that ActiveClean makes more progress towards the clean model with a smaller number of examples cleaned. To achieve a 2% error for the block removal, ActiveClean can inspect 2200 fewer images than Active Learning and 2750 fewer images than SampleClean. For the fuzzy images, both Active Learning and ActiveClean reach 2% error after cleaning fewer than 100 images, while SampleClean requires 1750.

Summary: In the MNIST dataset, ActiveClean significantly reduces (more than 2x) the number of records to clean to train a model with 2% error.




Figure 12: In a real adaptive detection scenario with the MNIST dataset, ActiveClean outperforms Active Learning and SampleClean.

9. RELATED WORK

Stochastic Optimization and Active Learning: Zhao and Tong recently proposed using importance sampling in conjunction with stochastic gradient descent [45]. However, calculating the optimal importance sampling distribution is very expensive and is only justified in this problem because data cleaning is even more expensive. Zhao and Tong use a different approximation to work around this problem. This work is one of many in an emerging consensus in stochastic optimization that not all data are equal (e.g., [34]). This line of work builds on prior results in linear algebra that show that some matrix columns are more informative than others [15], and Active Learning which shows that some labels are more informative than others [36]. Active Learning largely studies the problem of label acquisition [36], and recently the links between Active Learning and Stochastic optimization have been studied [20]. We use the work in Guillory et al. to evaluate a state-of-the-art Active Learning technique against ActiveClean.

Transfer Learning and Bias Mitigation: ActiveClean has a strong link to a field called Transfer Learning and Domain Adaptation [32]. The basic idea of Transfer Learning is that suppose a model is trained on a dataset D but tested on a dataset D' . Much of the complexity and contribution of this work comes from efficiently tuning such a process for expensive data cleaning applications – costs not studied in this field. In robotics, Mahler et al. explored a calibration problem in which data was systematically corrupted [28] and proposed a rule-based technique for cleaning data. Other problems in bias mitigation (e.g., Krishnan et al. [26]) have the same structure, systematically corrupted data that is feeding into a model. In this work, we try to generalize these principles given a general dirty dataset, convex model, and data cleaning procedure.

Secure Learning: Another relevant line of work is the work in private machine learning [16,41]. Learning is performed on a noisy variant of the data which mitigates privacy concerns. The goal is to extrapolate the insights from the noisy data to the hidden real data. ActiveClean is also related work in adversarial learning [30], where the goal is to make models robust to adversarial data manipulation. This line of work has extensively studied methodologies for making models private to external queries and robust to malicious labels [44], but not the data cleaning problem.

Data Cleaning and Databases: There are also several recent results in data cleaning that we would like to highlight. Progressive data cleaning methodologies have been proposed, however, these techniques tend to be application agnostic [29]. Altowim et al. proposed a framework for progressive entity resolution [7]. Volkovs et al. explored a related topic of maintaining data cleaning rules that change over time [40]. Recently, in works such as SampleClean [43], the application (i.e., queries) are used to inform data cleaning methodology. When the workload is made up of aggregate queries, cleaning samples of data may suffice. Similarly, Bergman et al. explore the problem of query-oriented data cleaning [8]. Given a query they clean data relevant to that query. Bergman et al. does not explore the Machine Learning applications studied in this work. Deshpande et al. studied data acquisition in sensor networks [14]. They explored value of information based prioritization of data acquisition for estimating aggregate queries of sensor readings. Similarly, Jeffery et al. [23] explored similar prioritization based on value of information. We see this work as pushing prioritization further down the pipeline to the end analytics. Finally, incremental optimization methods like SGD have a connection to incremental materialized view maintenance as the argument for incremental maintenance over recomputation is similar (i.e., relatively sparse updates). Krishnan et al. explored how samples of materialized views can be maintained similar to how models are updated with a sample of clean data in this work [27]

10. DISCUSSION AND FUTURE WORK

The experimental results suggest the following conclusions about ActiveClean: (1) when the data corruption rate is relatively small (e.g., 5%), ActiveClean cleans fewer records than Active Learning or SampleClean to achieve the same model accuracy, (2) all of the optimizations in ActiveClean (importance sampling, detection, and estimation) lead to significantly more accurate models at small sample sizes, (3) only when corruption rates are very severe (e.g. 50%) , SampleClean out performs ActiveClean, and (4) two real-world scenarios demonstrate similar accuracy improvements where ActiveClean returns significantly more accurate models than SampleClean or Active Learning for the same number of records cleaned.

There are also a few additional points for discussion. ActiveClean provides guarantees for training error on models trained with progressive data cleaning, however, there are no such guarantees on test error. This work focuses on the problem where an analyst has a large amount of dirty data and would like explore data cleaning and predictive models on this dataset. By providing the analyst more accurate model estimates, the value of different data cleaning techniques can be judged without having to clean the entire dataset. However, this problem is distinct from the model deployment problem, which we hope to explore in more detail in future work. It implicitly assumes that when the model is deployed, it will be applied in a setting where the test data is also clean. Training on clean data, and testing on dirty data, defeats the purpose of data cleaning and can lead to unreliable predictions.

As the experiments clearly show, ActiveClean is not strictly *better* than Active Learning or Sample Clean. ActiveClean is optimized for a specific design point of sparse errors and small sample sizes, and the empirical results suggest it returns more accurate models in this setting. As sample sizes and error rates increase, the benefits of ActiveClean are reduced. Another consideration for future work is automatically selecting alternative techniques when ActiveClean is expected to perform poorly.

Beyond these limitations, there are several exciting new avenues for future work. The data cleaning models explored in this work can be extended to handle non-uniform costs, where different errors have a different cleaning cost. Next, the empirical success of Deep Learning has led to increasing industry and research adoption of non-convex losses in many tasks that were traditionally served by convex models. In future work, we hope to explore how we can integrate with such frameworks.

11. CONCLUSION

The growing popularity of predictive models in data analytics adds additional challenges in managing dirty data. Progressive data cleaning in this setting is susceptible to errors due to mixing dirty and clean data, sensitivity to sample size, and the sparsity of errors. The key insight of ActiveClean is that an important class of predictive models, called convex loss models (e.g., linear regression and SVMs), can be simultaneously trained and cleaned. Consequently, there are provable guarantees on the convergence and error bounds of ActiveClean. ActiveClean also includes numerous optimizations such as: using the information from the model to inform data cleaning on samples, dirty data detection to avoid sampling clean data, and batching together updates on already clean data. The experimental results are promising as they suggest that these optimizations can significantly reduce data cleaning costs when errors are sparse and cleaning budgets are small. Techniques such as Active Learning and SampleClean are not optimized for the sparse low-budget setting, and ActiveClean achieves models of similar accuracy for significantly less records cleaned.

12. REFERENCES

- [1] Berkeley data analytics stack.
<https://amplab.cs.berkeley.edu/software/>.
- [2] Dollars for docs. <http://projects.propublica.org/open-payments/>.
- [3] For big-data scientists, 'janitor work' is key hurdle to insights.
<http://www.nytimes.com/2014/08/18/technology/for-big-data-scientists-hurdle-to-insights-is-janitor-work.html>.
- [4] A pharma payment a day keeps docs' finances okay.
<https://www.propublica.org/article/a-pharma-payment-a-day-keeps-docs-finances-ok>.
- [5] Sampleclean. <http://sampleclean.org/>, 2015.
- [6] A. Alexandrov, R. Bergmann, S. Ewen, J. Freytag, F. Hueske, A. Heise, O. Kao, M. Leich, U. Leser, V. Markl, F. Naumann, M. Peters, A. Rheinländer, M. J. Sax, S. Schelter, M. Höger, K. Tzoumas, and D. Warneke. The stratosphere platform for big data analytics. *VLDB J.*, 23(6):939–964, 2014.
- [7] Y. Altowim, D. V. Kalashnikov, and S. Mehrotra. Progressive approach to relational entity resolution. *Proceedings of the VLDB Endowment*, 7(11), 2014.
- [8] M. Bergman, T. Milo, S. Novgorodov, and W.-C. Tan. Query-oriented data cleaning with oracles. 2015.
- [9] D. P. Bertsekas. Incremental gradient, subgradient, and proximal methods for convex optimization: A survey. *Optimization for Machine Learning*, 2010:1–38.
- [10] L. Bottou. Stochastic gradient descent tricks. In *Neural Networks: Tricks of the Trade*, pages 421–436. Springer, 2012.
- [11] X. Chu, J. Morcos, I. F. Ilyas, M. Ouzzani, P. Papotti, N. Tang, and Y. Ye. KATARA: A data cleaning system powered by knowledge bases and crowdsourcing. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data, Melbourne, Victoria, Australia, May 31 - June 4, 2015*, pages 1247–1261, 2015.
- [12] A. Crotty, A. Galakatos, and T. Kraska. Tupleware: Distributed machine learning on small clusters. *IEEE Data Eng. Bull.*, 37(3):63–76, 2014.
- [13] O. Dekel, R. Gilad-Bachrach, O. Shamir, and L. Xiao. Optimal distributed online prediction using mini-batches. *The Journal of Machine Learning Research*, 13(1):165–202, 2012.
- [14] A. Deshpande, C. Guestrin, S. Madden, J. M. Hellerstein, and W. Hong. Model-driven data acquisition in sensor networks. In *(e)Proceedings of the Thirtieth International Conference on Very Large Data Bases, Toronto, Canada, August 31 - September 3 2004*, pages 588–599, 2004.
- [15] P. Drineas, M. Magdon-Ismail, M. W. Mahoney, and D. P. Woodruff. Fast approximation of matrix coherence and statistical leverage. *The Journal of Machine Learning Research*, 13(1):3475–3506, 2012.
- [16] J. C. Duchi, M. I. Jordan, and M. J. Wainwright. Local privacy and statistical minimax rates. In *Foundations of Computer Science (FOCS), 2013 IEEE 54th Annual Symposium on*, pages 429–438. IEEE, 2013.
- [17] J. Feng, H. Xu, S. Mannor, and S. Yan. Robust logistic regression and classification. In *Advances in Neural Information Processing Systems*, pages 253–261, 2014.
- [18] J. Friedman, T. Hastie, and R. Tibshirani. *The elements of statistical learning*, volume 1. Springer series in statistics Springer, Berlin, 2001.
- [19] C. Gokhale, S. Das, A. Doan, J. F. Naughton, N. Rampalli, J. Shavlik, and X. Zhu. Corleone: Hands-off crowdsourcing for entity matching. In *SIGMOD*, 2014.
- [20] A. Guillory, E. Chastain, and J. Bilmes. Active learning as non-convex optimization. In *International Conference on Artificial Intelligence and Statistics*, pages 201–208, 2009.
- [21] J. M. Hellerstein, C. Ré, F. Schoppmann, D. Z. Wang, E. Fratkin, A. Gorajek, K. S. Ng, C. Welton, X. Feng, K. Li, and A. Kumar. The madlib analytics library or MAD skills, the SQL. *PVLDB*, 5(12):1700–1711, 2012.
- [22] M. Jaggi, V. Smith, M. Takác, J. Terhorst, S. Krishnan, T. Hofmann, and M. I. Jordan. Communication-efficient distributed dual coordinate ascent. In *Advances in Neural Information Processing Systems*, pages 3068–3076, 2014.
- [23] S. R. Jeffery, G. Alonso, M. J. Franklin, W. Hong, and J. Widom. Declarative support for sensor data cleaning. In *Pervasive*, pages 83–100, 2006.
- [24] Z. Khayyat, I. F. Ilyas, A. Jindal, S. Madden, M. Ouzzani, P. Papotti, J.-A. Quiané-Ruiz, N. Tang, and S. Yin. Bigdancing: A system for big data cleansing. 2015.
- [25] H. Köpcke, A. Thor, and E. Rahm. Evaluation of entity resolution approaches on real-world match problems. *PVLDB*, 3(1):484–493, 2010.
- [26] S. Krishnan, J. Patel, M. J. Franklin, and K. Goldberg. A methodology for learning, analyzing, and mitigating social influence bias in recommender systems. In *Eighth ACM Conference on Recommender Systems, RecSys '14, Foster City, Silicon Valley, CA, USA - October 06 - 10, 2014*, pages 137–144, 2014.
- [27] S. Krishnan, J. Wang, M. J. Franklin, K. Goldberg, and T. Kraska. Stale view cleaning: Getting fresh answers from stale materialized views. *Proceedings of the VLDB Endowment*, 8(12), 2015.
- [28] J. Mahler, S. Krishnan, M. Laskey, S. Sen, A. Murali, B. Kehoe, S. Patil, J. Wang, M. Franklin, P. Abbeel, and K. Y. Goldberg. Learning accurate kinematic control of cable-driven surgical robots using data cleaning and gaussian process regression. In *2014 IEEE International Conference on Automation Science and Engineering, CASE 2014, New Taipei, Taiwan, August 18-22, 2014*, pages 532–539, 2014.
- [29] C. Mayfield, J. Neville, and S. Prabhakar. ERACER: a database approach for statistical inference and data cleaning. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2010, Indianapolis, Indiana, USA, June 6-10, 2010*, pages 75–86, 2010.
- [30] B. Nelson, B. I. Rubinstein, L. Huang, A. D. Joseph, S. J. Lee, S. Rao, and J. Tygar. Query strategies for evading convex-inducing classifiers. *The Journal of Machine Learning Research*, 13(1):1293–1332, 2012.
- [31] A. B. Owen. *Monte Carlo theory, methods and examples*. 2013.
- [32] S. J. Pan and Q. Yang. A survey on transfer learning. *Knowledge and Data Engineering, IEEE Transactions on*, 22(10):1345–1359, 2010.
- [33] H. Park and J. Widom. Crowdfill: collecting structured data from the crowd. In *International Conference on Management of Data, SIGMOD 2014, Snowbird, UT, USA, June 22-27, 2014*, pages 577–588, 2014.
- [34] Z. Qu, P. Richtárik, and T. Zhang. Randomized dual coordinate ascent with arbitrary sampling. *arXiv preprint arXiv:1411.5873*, 2014.
- [35] E. Rahm and H. H. Do. Data cleaning: Problems and current approaches. *IEEE Data Eng. Bull.*, 23(4):3–13, 2000.
- [36] B. Settles. Active learning literature survey. *University of Wisconsin, Madison*, 52(55-66):11, 2010.
- [37] E. H. Simpson. The interpretation of interaction in contingency tables. *Journal of the Royal Statistical Society. Series B (Methodological)*, pages 238–241, 1951.
- [38] N. Swartz. Gartner warns firms of 'dirty data'. *Information Management Journal*, 41(3), 2007.
- [39] J. R. Taylor. An introduction to error analysis: The study of uncertainties in physical measurements, 327 pp. *Univ. Sci. Books, Mill Valley, Calif.*, 1982.
- [40] M. Volkovs, F. Chiang, J. Szlichta, and R. J. Miller. Continuous data cleaning. In *Data Engineering (ICDE), 2014 IEEE 30th International Conference on*, pages 244–255. IEEE, 2014.
- [41] M. J. Wainwright, M. I. Jordan, and J. C. Duchi. Privacy aware learning. In *Advances in Neural Information Processing Systems*, pages 1430–1438, 2012.
- [42] J. Wang, T. Kraska, M. J. Franklin, and J. Feng. Crowder: Crowdsourcing entity resolution. *PVLDB*, 5(11):1483–1494, 2012.
- [43] J. Wang, S. Krishnan, M. J. Franklin, K. Goldberg, T. Kraska, and T. Milo. A sample-and-clean framework for fast and accurate query processing on dirty data. In *SIGMOD Conference*, pages 469–480, 2014.
- [44] H. Xiao, T. DE, B. Biggio, D. UNICA, G. Brown, G. Fumera, C. Eckert, I. TUM, and F. Roli. Is feature selection secure against training data poisoning? 2015.
- [45] P. Zhao and T. Zhang. Stochastic optimization with importance sampling. *arXiv preprint arXiv:1401.2753*, 2014.

APPENDIX

A. EXTENSIONS

A.1 Set-of-Records Cleaning Model

In paper, we formalized the analyst-specified data cleaning as follows. We take the sample of the records S_{dirty} , and apply data cleaning $C(\cdot)$. C is applied to a record and produces the clean record:

$$S_{clean} = \{C(r) : \forall r \in S_{dirty}\}$$

The record-by-record cleaning model is a formalization of the costs of data cleaning where each record has the same cost to clean and this cost does not change throughout the entire cleaning session. There are, however, some cases when cleaning the first record of a certain type of corruption is expensive but all subsequent records are cheaper.

EXAMPLE 6. *In most spell checking systems, when a misspelling is identified, the system gives an option to fix all instances of that misspelling.*

EXAMPLE 7. *In entity resolution problems, when an inconsistent entity is identified all other records with the same inconsistency can be efficiently fixed.*

This model of data cleaning can fit into our framework and we formalize it as the ‘‘Set-of-Records’’ model as opposed to the ‘‘Record-by-Record’’ model. In this model, the cleaning function $C(\cdot)$ is not restricted to updating only the records in the sample. $C(\cdot)$ takes the entire dirty sample as an argument (that is the cleaning is a function of the sample), the dirty data, and updates the entire dirty data:

$$R'_{dirty} = C(S_{dirty}, R_{dirty})$$

we require that for every record $s \in S_{dirty}$, that record is completely cleaned after applying $C(\cdot)$, giving us S_{clean} . Records outside of S_{dirty} may be cleaned on a subset of dirty attributes by $C(\cdot)$. After each iteration, we re-run the detector, and move any $r \in R'_{dirty}$ that are clean to R_{clean} . Such a model allows us to capture data cleaning operations such as in Example 6 and Example 7.

A.2 Non-convex losses

We acknowledge that there is an increasing popularity of non-convex losses in the Neural Network and Deep Learning literature. However, even for these losses, gradient descent techniques still apply. Instead of converging to a global optimum they converge to a locally optimal value. Likewise, ActiveClean will converge to the closest locally optimal value to the dirty model. Because of this, it is harder to reason about the results. Different initializations will lead to different local optima, and thus, introduces a complex dependence on the initialization with the dirty model. This problem is not fundamental to ActiveClean and any gradient technique suffers this challenge for general non-convex losses, and we hope to explore this more in the future.

B. PROOF OF LEMMA 1

LEMMA 5. *The gradient estimate $g(\theta)$ is unbiased if g_S is an unbiased estimate of:*

$$\frac{1}{|R_{dirty}|} \sum g_i(\theta)$$

PROOF SKETCH.

$$\mathbb{E}\left(\frac{1}{|R_{dirty}|} \sum g_i(\theta)\right) = \frac{1}{|R_{dirty}|} \cdot \mathbb{E}\left(\sum g_i(\theta)\right)$$

By symmetry,

$$\mathbb{E}\left(\frac{1}{|R_{dirty}|} \sum g_i(\theta)\right) = g(\theta)$$

$$\mathbb{E}\left(\frac{1}{|R_{dirty}|} \sum g_i(\theta)\right) = \frac{|R_{dirty}| \cdot g_S + |R_{clean}| \cdot g_C}{|R|}$$

□

C. CONVEX VS. STRONGLY CONVEX

The error bound discussed in Proposition 2 can be tightened for a class of models called strongly convex (see [9] for a definition). Linear regression, logistic regression, and regularized support vector machines fit into this class.

PROPOSITION 3. *For a strongly convex loss, a batch size b , and T iterations, the convergence rate is bounded by $O(\frac{\sigma^2}{bT})$.*

D. PROOF OF LEMMA 2

The variance of this estimate is given by:

$$Var(\mu) = \mathbb{E}(\mu^2) - \mathbb{E}(\mu)^2$$

Since the estimate is unbiased, we can replace $\mathbb{E}(\mu)$ with the average of A :

$$Var(\mu) = \mathbb{E}(\mu^2) - \bar{A}^2$$

Since \bar{A} is deterministic, we can remove that term during minimization. Furthermore, we can write $\mathbb{E}(\mu^2)$ as:

$$\mathbb{E}(\mu^2) = \frac{1}{n^2} \sum_i \frac{a_i^2}{p_i}$$

Then, we can solve the following optimization problem (removing the proportionality of $\frac{1}{n^2}$) over the set of weights $P = \{p(a_i)\}$:

$$\min_P \sum_i \frac{a_i^2}{p_i}$$

$$\text{subject to: } P > 0, \sum P = 1$$

Applying Lagrange multipliers, an equivalent unconstrained optimization problem is:

$$\min_{P>0, \lambda>0} \sum_i \frac{a_i^2}{p_i} + \lambda \cdot (\sum P - 1)$$

If, we take the derivatives with respect to p_i and set them equal to zero:

$$-\frac{a_i^2}{2 \cdot p_i^2} + \lambda = 0$$

If, we take the derivative with respect to λ and set it equal to zero:

$$\sum P - 1$$

Solving the system of equations, we get:

$$p_i = \frac{|a_i|}{\sum_i |a_i|}$$

E. TAYLOR APPROXIMATION

We can take the expected value of the Taylor series expansion around the dirty value. If d is the dirty value and c is the clean value, the Taylor series approximation for a function f is given as follows:

$$f(c) = f(d) + f'(d) \cdot (d - c) + \dots$$

If we ignore the higher order terms, we see that the linear term $f'(d) \cdot (d - c)$ decouples the features. We only have to know the change in each feature to estimate the change in value. In our case the function f is the gradient $\nabla\phi$. So, the resulting linearization is:

$$\begin{aligned} \nabla\phi(x_i^{(c)}, y_i^{(c)}, \theta) &\approx \nabla\phi(x, y, \theta) + \frac{\partial}{\partial X} \nabla\phi(x, y, \theta) \cdot (x - x^{(c)}) \\ &\quad + \frac{\partial}{\partial Y} \nabla\phi(x, y, \theta) \cdot (y - y^{(c)}) \end{aligned}$$

When we take the expected value:

$$\begin{aligned} \mathbb{E}(\nabla\phi(x_{clean}, y_{clean}, \theta)) &\approx \nabla\phi(x, y, \theta) + \frac{\partial}{\partial X} \nabla\phi(x, y, \theta) \cdot \mathbb{E}(\Delta x) \\ &\quad + \frac{\partial}{\partial Y} \nabla\phi(x, y, \theta) \cdot \mathbb{E}(\Delta y) \end{aligned}$$

So the resulting estimation formula takes the following form:

$$\approx \nabla\phi(x, y, \theta) + M_x \cdot \mathbb{E}(\Delta x) + M_y \cdot \mathbb{E}(\Delta y)$$

Recall that we have a d dimensional feature space and l dimensional label space. Then, $M_x = \frac{\partial}{\partial X} \nabla\phi$ is an $d \times d$ matrix, and $M_y = \frac{\partial}{\partial Y} \nabla\phi$ is a $d \times l$ matrix. Both of these matrices are computed with respect to dirty data, and we will present an example. Δx is a d dimensional vector where each component represents a change in that feature and Δy is an l dimensional vector that represents the change in each of the labels.

F. EXAMPLE M_X, M_Y

Linear Regression:

$$\nabla\phi(x, y, \theta) = (\theta^T x - y)x$$

For a record, r , suppose we have a feature vector x . If we take the partial derivatives with respect to x , M_x is:

$$M_x[i, i] = 2x[i] + \sum_{i \neq j} \theta[j]x[j] - y$$

$$M_x[i, j] = \theta[j]x[i]$$

Similarly M_y is:

$$M_y[i, 1] = x[i]$$

Logistic Regression:

$$\nabla\phi(x, y, \theta) = (h(\theta^T x) - y)x$$

where

$$h(z) = \frac{1}{1 + e^{-z}}$$

we can rewrite this as:

$$h_\theta(x) = \frac{1}{1 + e^{\theta^T x}}$$

$$\nabla\phi(x, y, \theta) = (h_\theta(x) - y)x$$

In component form,

$$g = \nabla\phi(x, y, \theta)$$

$$g[i] = h_\theta(x) \cdot x[i] - yx[i]$$

Therefore,

$$M_x[i, i] = h_\theta(x) \cdot (1 - h_\theta(x)) \cdot \theta[i]x[i] + h_\theta(x) - y$$

$$M_x[i, j] = h_\theta(x) \cdot (1 - h_\theta(x)) \cdot \theta[j]x[i] + h_\theta(x)$$

$$M_y[i, 1] = x[i]$$

SVM:

$$\nabla\phi(x, y, \theta) = \begin{cases} -y \cdot x & \text{if } y \cdot x \cdot \theta \leq 1 \\ 0 & \text{if } y \cdot x \cdot \theta \geq 1 \end{cases}$$

Therefore,

$$M_x[i, i] = \begin{cases} -y[i] & \text{if } y \cdot x \cdot \theta \leq 1 \\ 0 & \text{if } y \cdot x \cdot \theta \geq 1 \end{cases}$$

$$M_x[i, j] = 0$$

$$M_y[i, 1] = x[i]$$

G. AGGREGATE QUERIES AS CONVEX LOSSES

G.1 AVG and SUM queries

avg, sum queries are a special case of the convex loss minimization discussed in the paper: If we define the following loss, it is easy to verify the the optimal θ is the mean μ :

$$\phi = (x_i - \theta)^2$$

with the appropriate scaling it can support avg, sum queries with and without predicates. Taking the gradient of that loss:

$$\nabla\phi = 2(x_i - \theta)$$

It is also easy to verify that the bound on errors is $O(\frac{\mathbb{E}((x-\mu)^2)}{bT})$, which is essentially the CLT. The importance sampling results are intuitive as well. Applying the linearization:

$$M_x = 2$$

The importance sampling prioritizes points that it expects to be far away from the mean.

G.2 MEDIAN

Similarly, we can analyze the median query. If we define the following loss, it is easy to verify the the optimal θ is the median m :

$$\phi = |x_i - \theta|$$

Taking the gradient of that loss:

$$\nabla\phi = 1 \text{ if } < m, -1 \text{ if } > m$$

Applying the linearization:

$$M_x = 0$$

The intuitive result is that a robust query like a median does not need to consider estimation as the query result is robust to small changes.

H. EXPERIMENTAL COMPARISON

H.1 Robust Logistic Regression

We use the algorithm from Feng et al. for robust logistic regression.

1. Input: Contaminated training samples $\{(x_1, y_1), \dots, (x_n, y_n)\}$ an upper bound on the number of outliers n , number of inliers n and sample dimension p .

2. Initialization: Set

$$T = 4\sqrt{\log p/n + \log n/n}$$

3. Remove samples (x_i, y_i) whose magnitude satisfies $\|x_i\| \geq T$.
4. Solve regularized logistic regression problem.

H.2 Active Learning

There is a well established link between Active Learning and Online Gradient-based Algorithms [20]. To fairly evaluate an Active Learning methodology in these experiments, we run a gradient descent where examples are prioritized by expected gradient length (explained in [36]). Ignoring the detection step, there are two differences between this algorithm and the one we propose. First, we calculate the gradient with respect to the an estimate of the clean data, and second we sample rather than using a deterministic ordering. While such Active Learning algorithms have been studied in the Learning Theory community, they have not been adopted in data cleaning or crowdsourcing research. Typical algorithms include uncertainty sampling, where a classifier prioritized data closest to the margin. However, algorithms such as uncertainty sampling focus on the narrow problem of label acquisition in hyperplane classifiers; a problem too narrow for application in this setting.

I. DOLLARS FOR DOCS

The dollars for docs dataset has the following schema:

```
Contribution(pi_speciality, drug_name, device_name,
corporation, amount, dispute, status)
```

To flag suspect donations, we used the `status` attribute. When the `status` was "covered" that means it was an allowed contribution under the researcher's declared protocol. When the `status` was "non-covered" that means it was a disallowed contribution under the researcher's declared protocol. The rest of the textual attributes were featurized with a bag-of-words model, and the numerical amount and dispute attributes were treated as numbers.

We cleaned the entire Dollars for Docs dataset upfront to be able to evaluate how different budgeted data cleaning strategies compare to cleaning the full data. To clean the dataset, we loaded the entire data 240089 records into Microsoft Excel. We identified four broad classes of errors:

Corporations are inconsistently represented: "Pfizer", "Pfizer Inc.", "Pfizer Incorporated".

Drugs are inconsistently represented: "TAXOTERE DOCETAXEL -PROSTATE CANCER" and "TAXOTERE"

Label of covered and not covered are not consistent: "No", "Yes", "N", "This study is not supported", "None", "Combination"

Research subject must be a drug OR a medical device and not both: "BIO FLU QPAN H7N9AS03 Vaccine" and "BIO FLU QPAN H7N9AS03 Device"

To fix these errors, we sorted by each column and merged entities that looked similar and removed inconsistencies as in the status labels. When there were ambiguities, we referred to the drug company's website and whitepapers. When possible, we used batch

data transformations, like find and replace (i.e. the Set-of-Records model). In all, 44234 records had some error and full data cleaning required about 2 days of efforts.

Once cleaned, in our experiment, we encoded the 4 problems as data quality constraints. To fix the constraints, we looked up the clean value in the dataset that we cleaned up front.

Rule 1: Matching dependency on corporation (Weighted Jaccard Similarity > 0.8).

Rule 2: Matching dependency on drug (Weighted Jaccard Similarity > 0.8).

Rule 3: Label must either be "covered" or "not covered".

Rule 4: Either drug or medical device should be null.

J. MNIST ERRORS

We include visualization of the errors that we generated for the MNIST experiment. We generated these errors in MATLAB by taking the grayscale version of the image (a 64×64 matrix) and corrupting them by block removal and fuzzifying.

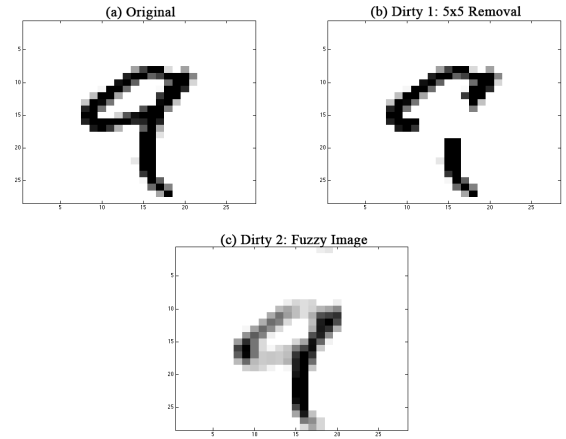


Figure 13: We experiment with two forms of corruption in the MNIST image datasets: 5x5 block removal and making the images fuzzy. Image (a) shows an uncorrupted "9", image (b) shows one corrupted with block removal, and image (c) shows one that is corrupted with fuzziness.